

# SAT/SMT Solver Basics

## (and Impact of Community Structure on SAT Solver Performance)

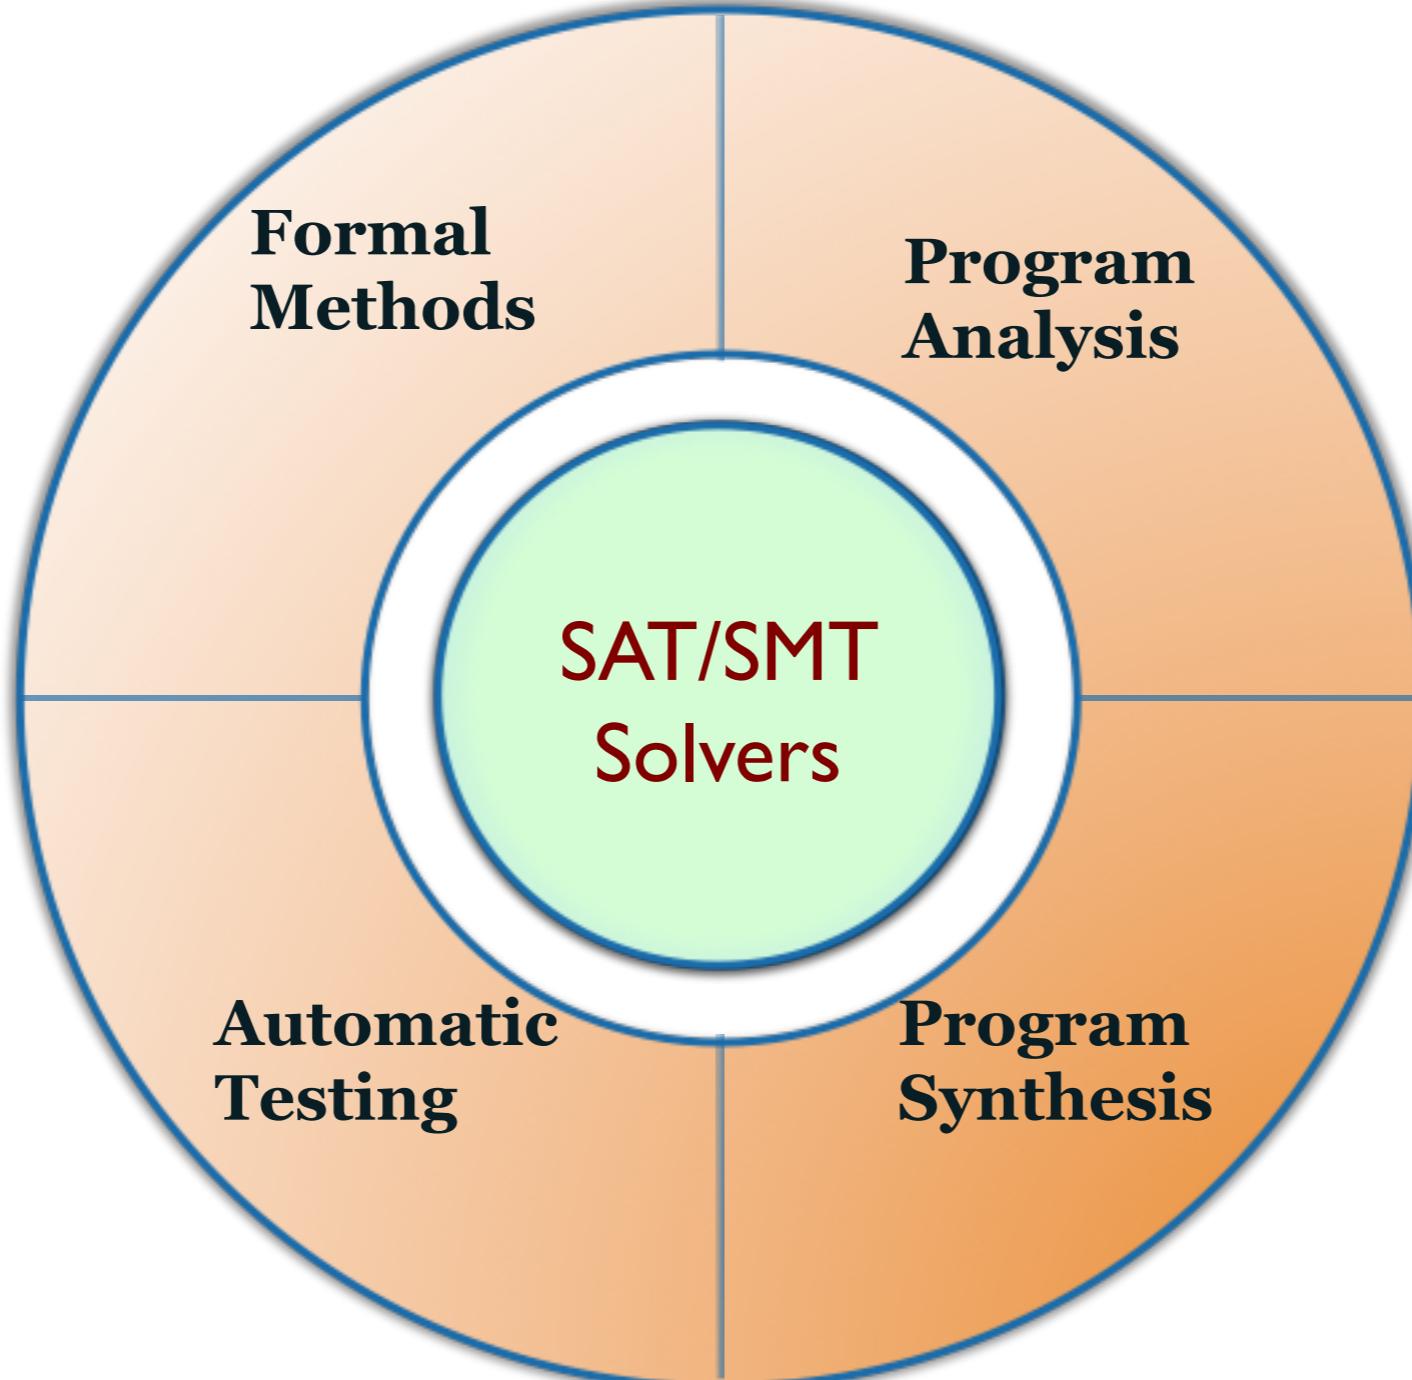
Vijay Ganesh

Affiliation: University of Waterloo

October 28, 2014, Dagstuhl, Germany

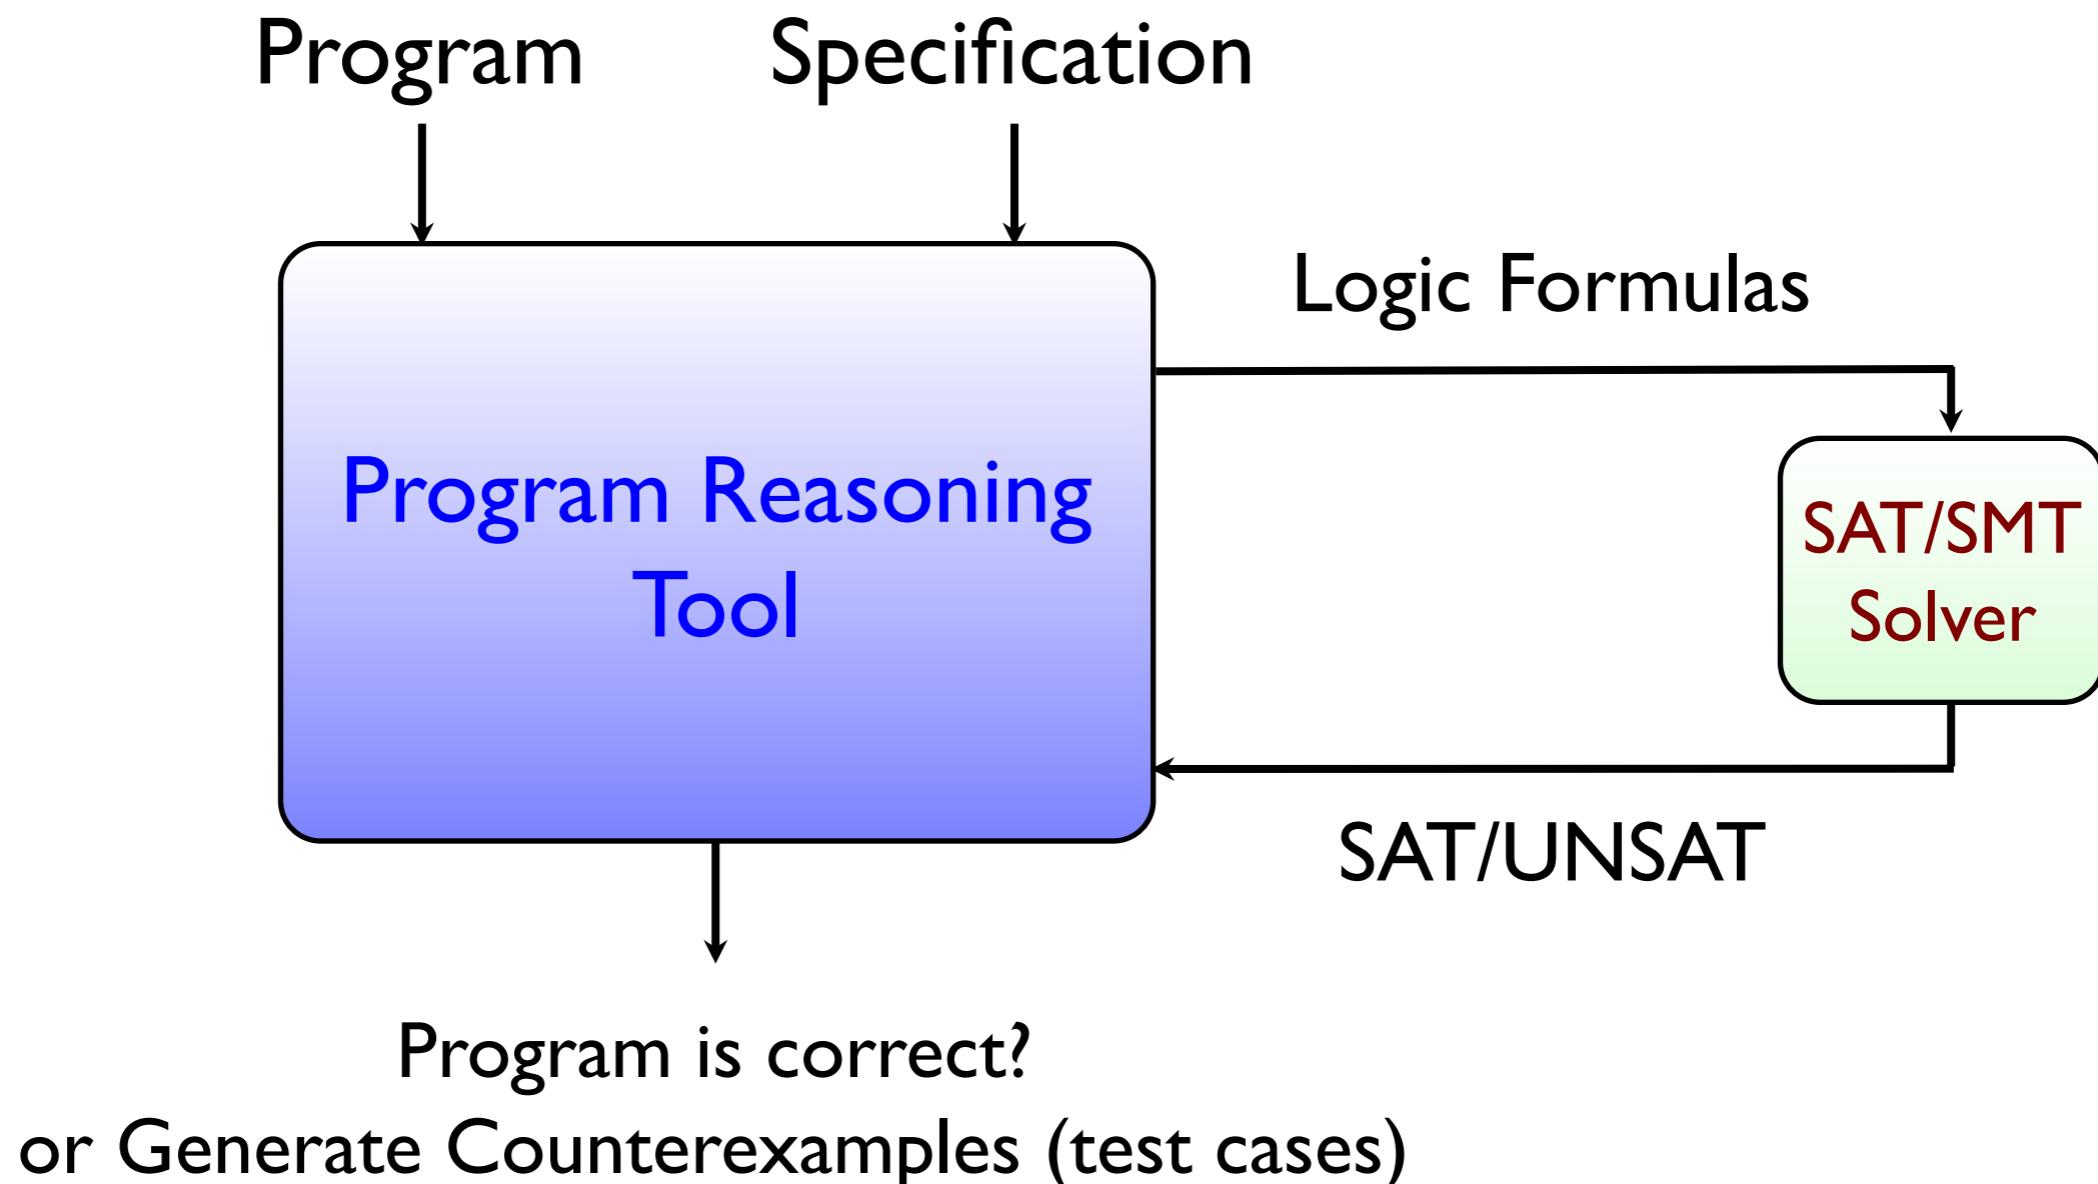
# Software Engineering & SAT/SMT Solvers

## An Indispensable Tactic for Any Strategy



# Software Engineering using Solvers

## Engineering, Usability, Novelty



# SAT/SMT Solver Research Story

## A 1000x Improvement

1,000,000 Constraints

- Solver-based programming languages
- Compiler optimizations using solvers
- Solver-based debuggers
- Solver-based type systems
- Solver-based concurrency bugfinding
- Solver-based synthesis
- Bio & Optimization

100,000 Constraints

- Concolic Testing
- Program Analysis
- Equivalence Checking
- Auto Configuration

10,000 Constraints

- Bounded MC
- Program Analysis
- AI

1,000 Constraints

1998

2000

2004

2007

2010

# Talk Outline

## **Topics covered on SAT Solvers**

- Motivation for SAT/SMT solvers in software engineering
- High-level description of the SAT problem
- Key techniques used in SAT solvers
  - DPLL search and its shortcomings
  - Modern CDCL SAT solver: propagate (BCP), decide (VSIDS), conflict analysis, clause learn and backJump, clause deletion
  - **Big lesson: learning from mistakes**

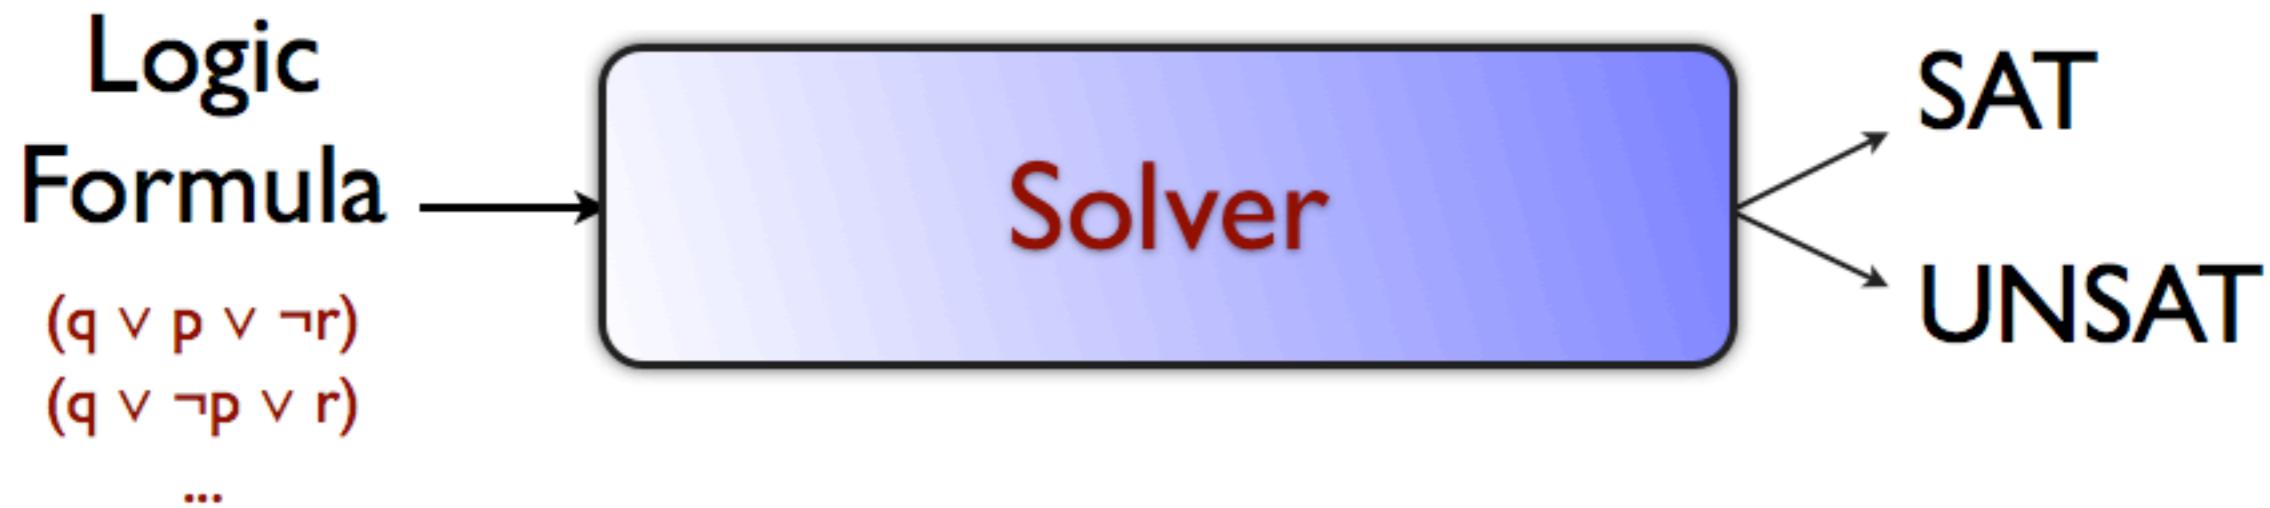
## **Topics covered on SMT Solvers**

- Modern SMT solvers and key techniques
  - Combination of theories
  - DPLL(T): Boolean reasoning over theories
  - Various forms of abstraction and refinement techniques
- Future directions

## **Impact of Community Structure on SAT Solver Performance (presented at SAT 2014)**

# What is a SAT/SMT Solver?

## Automation of Mathematical Logic



- Rich logics (Modular arithmetic, arrays, strings, non-linear arithmetic, theories with quantifiers, ...)
- From proof procedures to validity to satisfiability
- SAT problem is NP-complete, PSPACE-complete, ...
- Practical, scalable, usable, automatic
- Enable novel software reliability approaches

# DPLL SAT Solver Architecture

## The Basic Solver

```
DPLL( $\Theta_{cnf}$ , assign) {
```

Propagate unit clauses;

```
if "conflict": return FALSE;
```

```
if "complete assign": return TRUE;
```

```
"pick decision variable x";
```

```
return
```

```
    DPLL( $\Theta_{cnf}$  | x=0, assign[x=0])  
    || DPLL( $\Theta_{cnf}$  | x=1, assign[x=1]);
```

```
}
```

### Key Steps in a DPLL SAT Solver

#### Propagate (Boolean Constant Propagation)

- Propagate inferences due to unit clauses
- Most of the solving “effort” goes into this step

#### Detect Conflict

- Conflict: partial assignment is not satisfying

#### Decide (Branch)

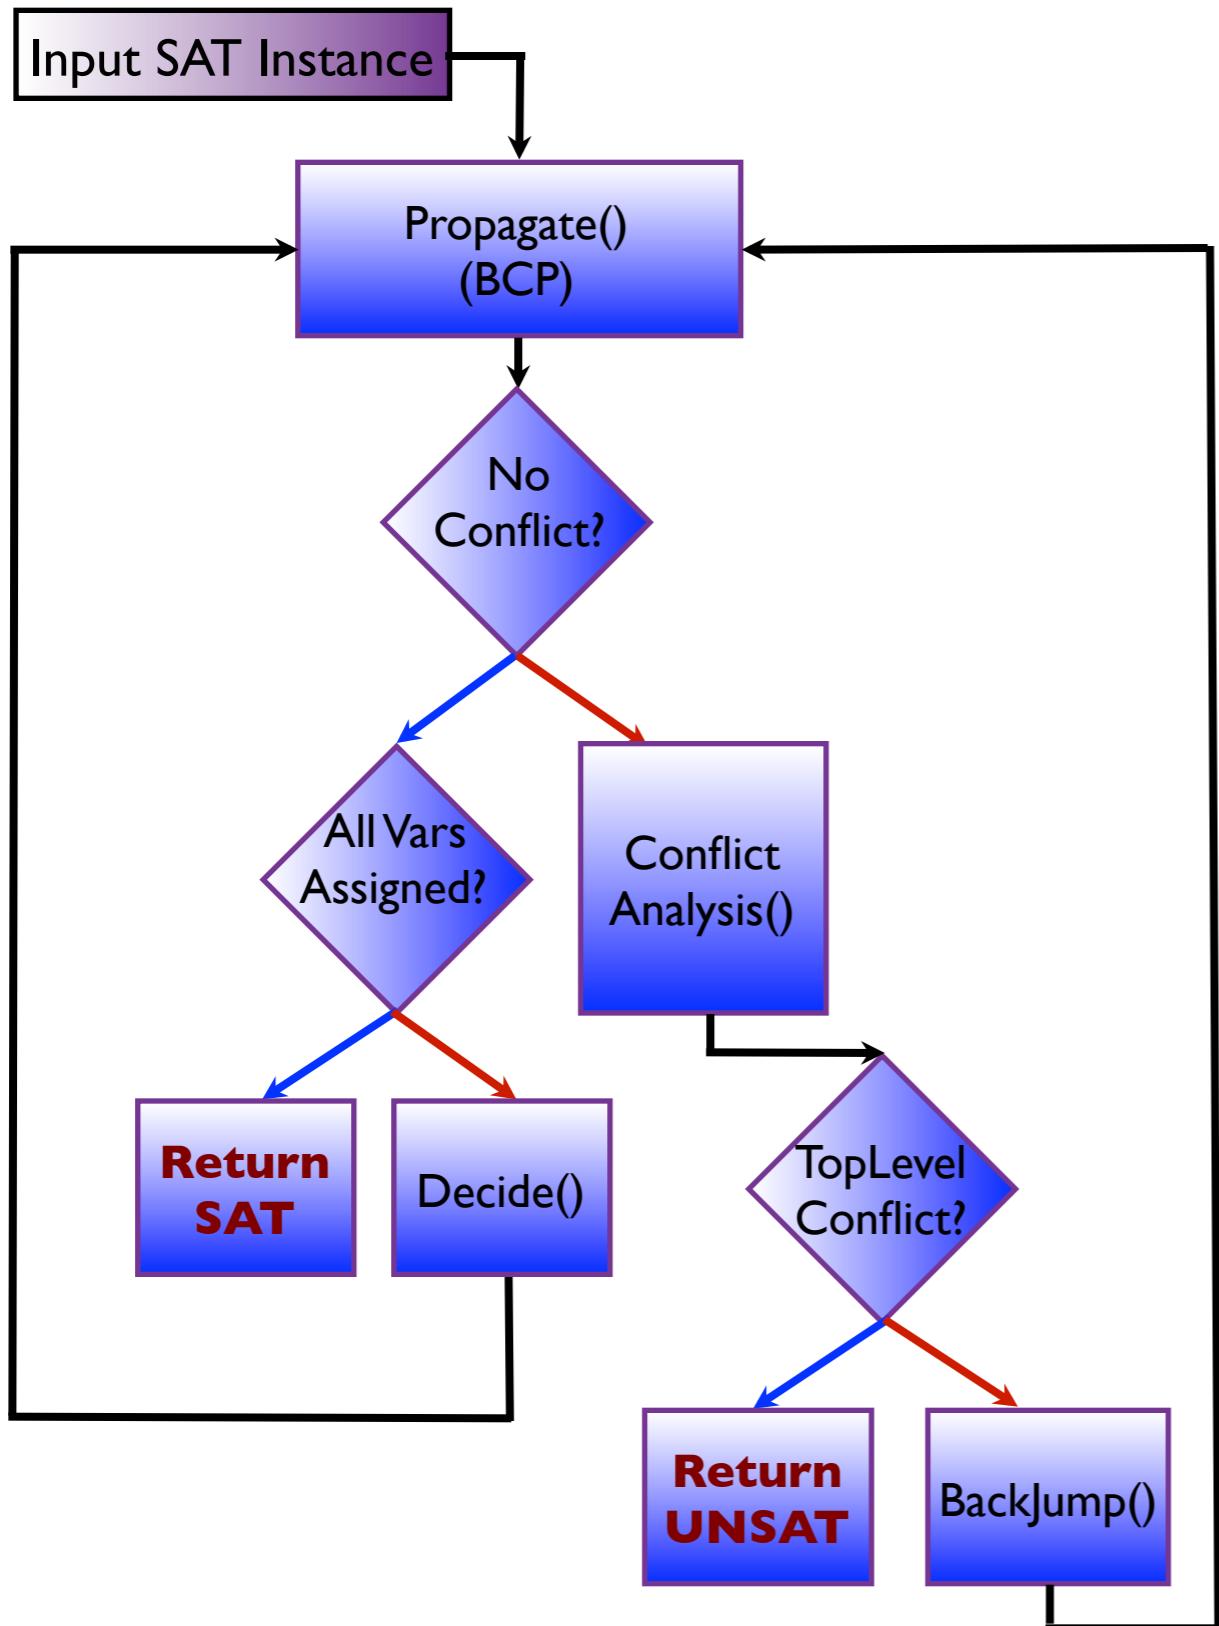
- Choose a variable & assign some value

#### Backtracking

- Implicitly done through the recursive call in DPLL

# Modern CDCL SAT Solver Architecture

## Key Steps and Data-structures



### Key steps

- Decide()
- Propagate() (Boolean constant propagation)
- Conflict analysis and learning() (CDCL)
- Backjump()
- Forget()
- Restart()

### CDCL: Conflict-Driven Clause-Learning

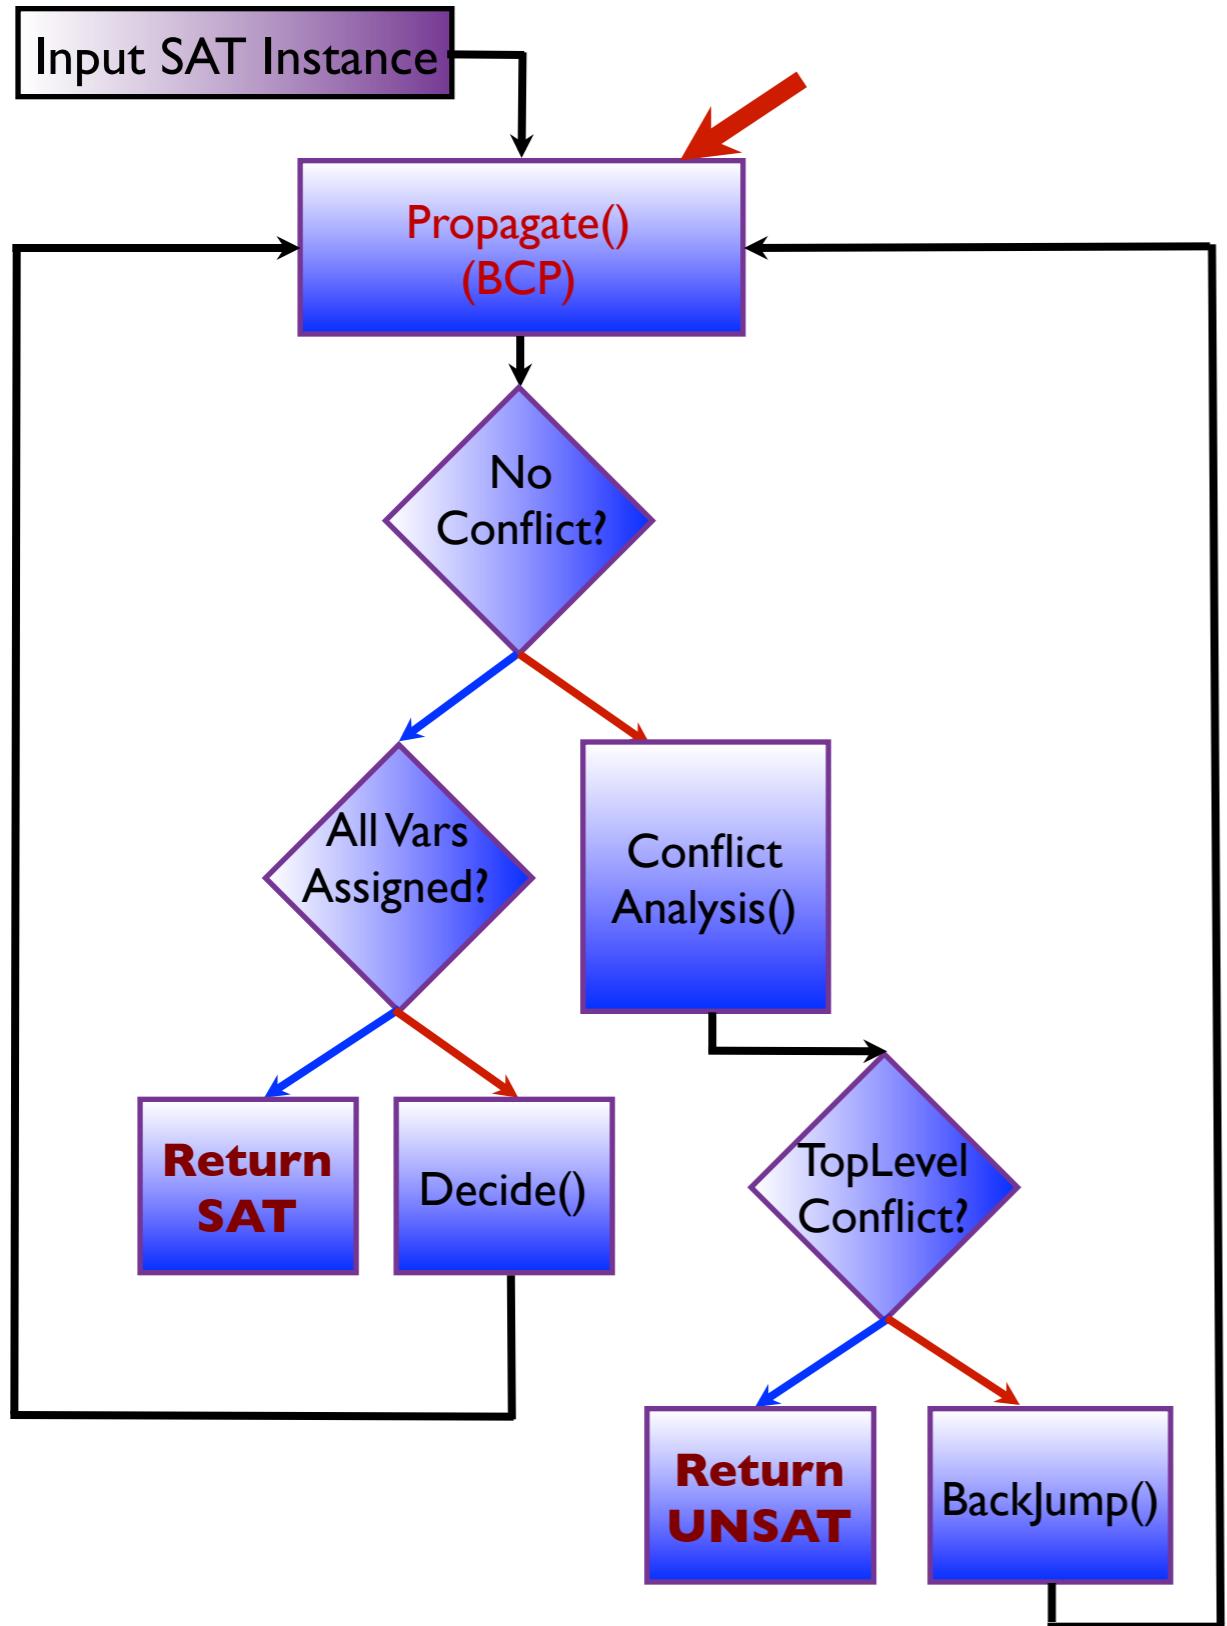
- Conflict analysis is a key step
- Results in learning a learnt clause
- Prunes the search space

### Key data-structures (Solver state)

- Stack or trail of partial assignments (AT)
- Input clause database
- Conflict clause database
- Conflict graph
- Decision level (DL) of a variable

# Modern CDCL SAT Solver Architecture

## Propagate(), Decide(), Analyze/Learn(), BackJump()

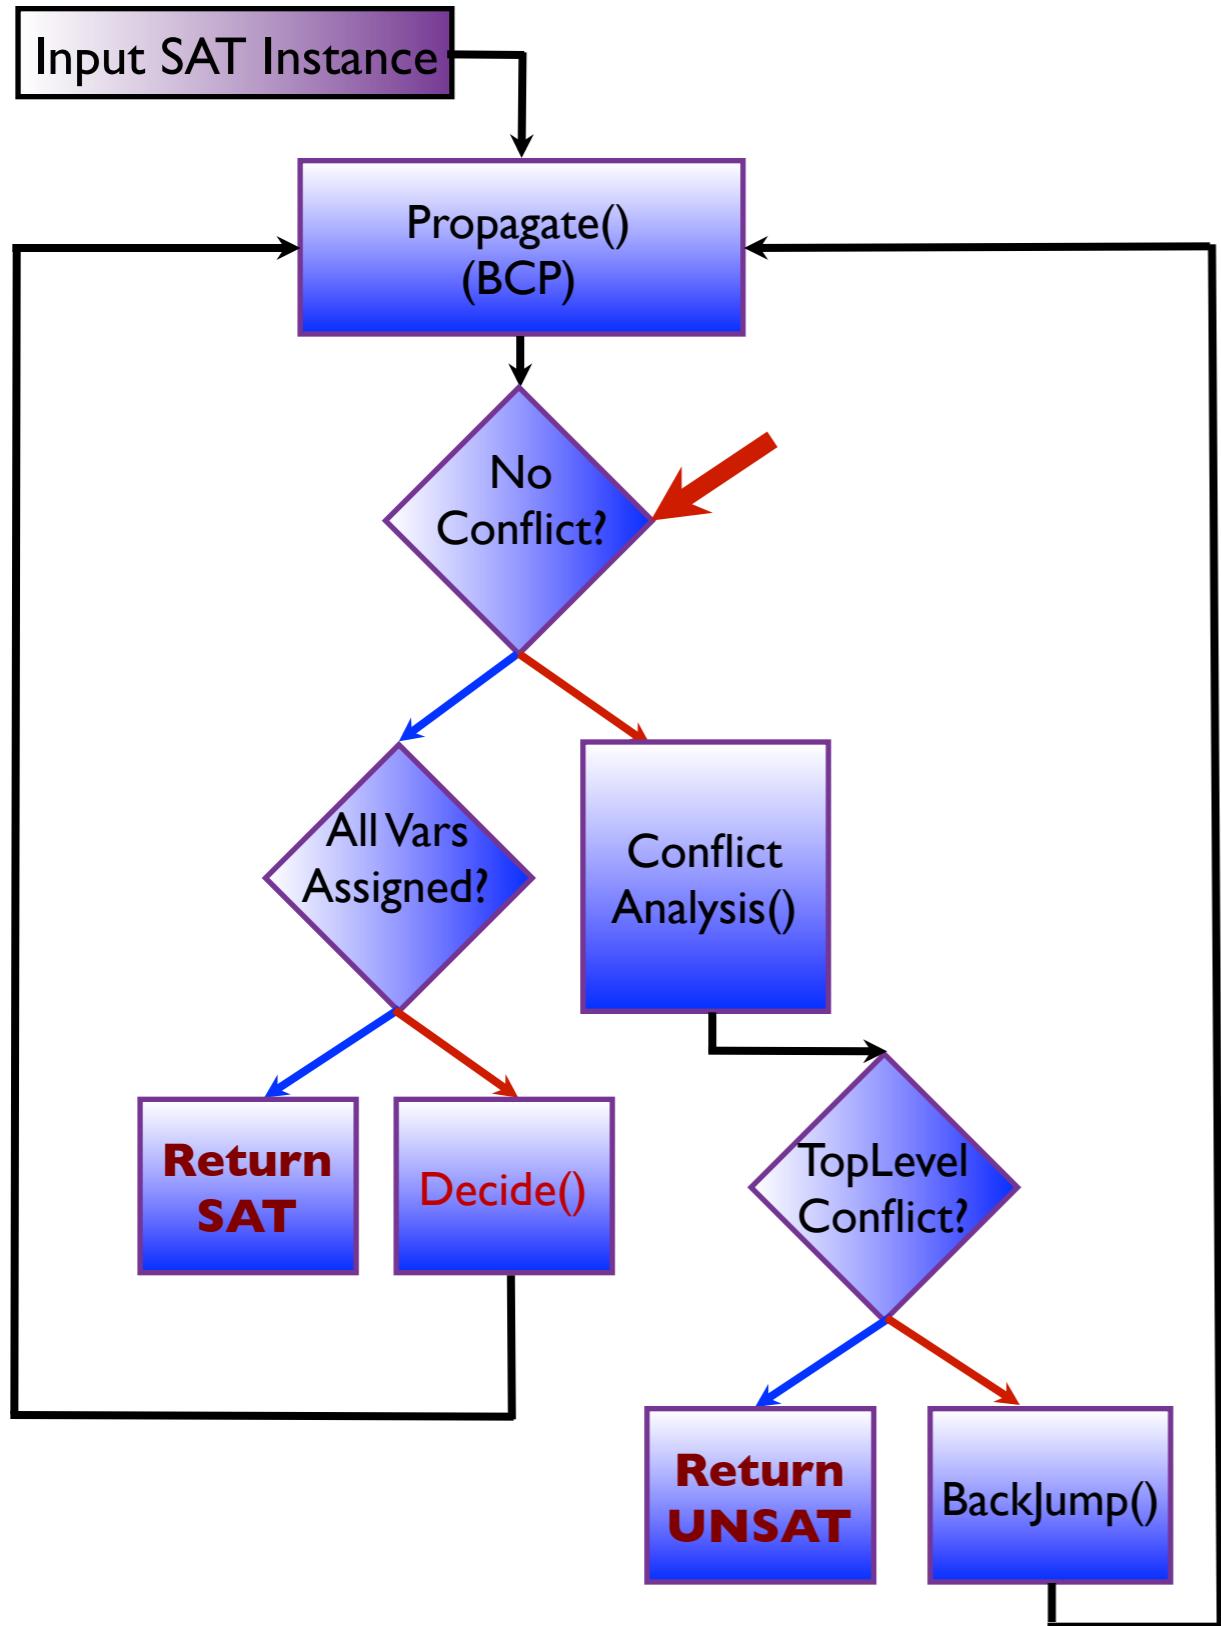


### Propagate (Boolean Constant Propagation)

- Propagate inferences due to unit clauses
- Most time in solving goes into this

# Modern CDCL SAT Solver Architecture

## Propagate(), Decide(), Analyze/Learn(), BackJump()



### Propagate (Boolean Constant Propagation)

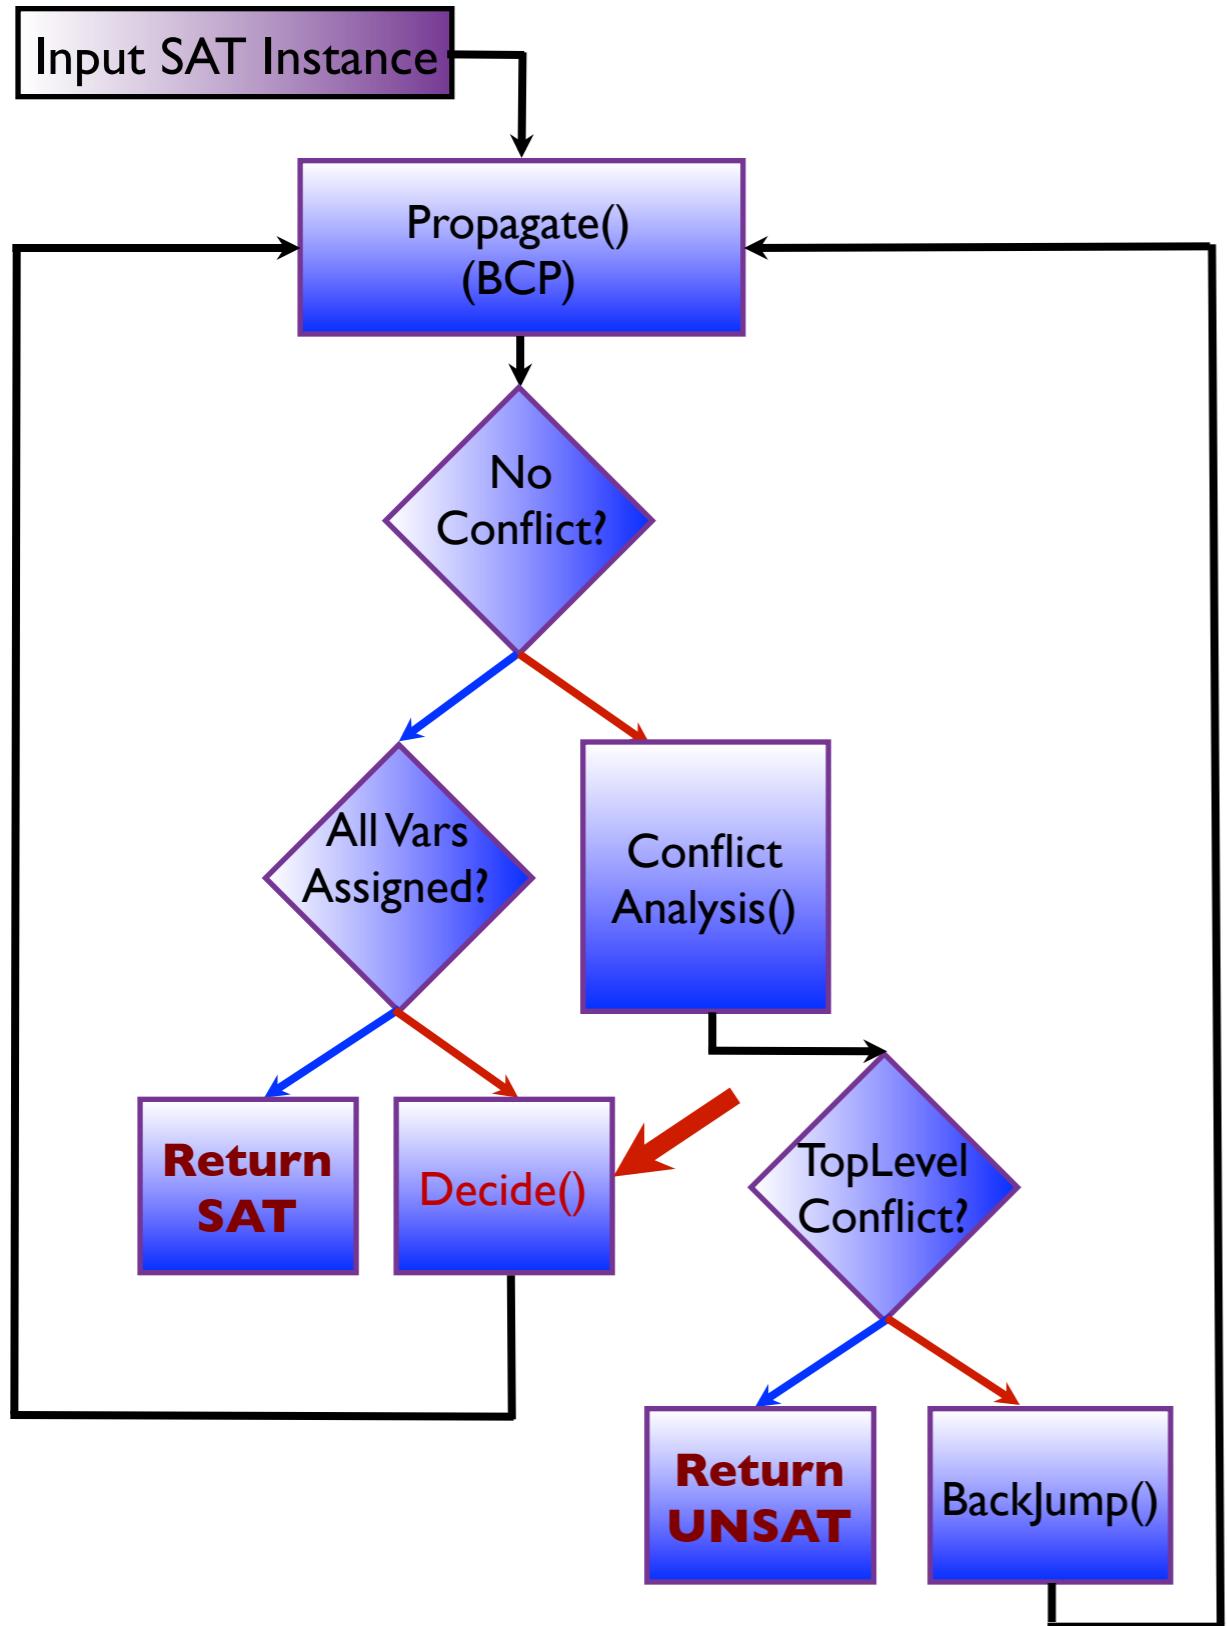
- Propagate inferences due to unit clauses
- Most time in solving goes into this

### Detect Conflict?

- Conflict: partial assignment is not satisfying

# Modern CDCL SAT Solver Architecture

## Propagate(), Decide(), Analyze/Learn(), BackJump()



### Propagate (Boolean Constant Propagation)

- Propagate inferences due to unit clauses
- Most time in solving goes into this

### Detect Conflict?

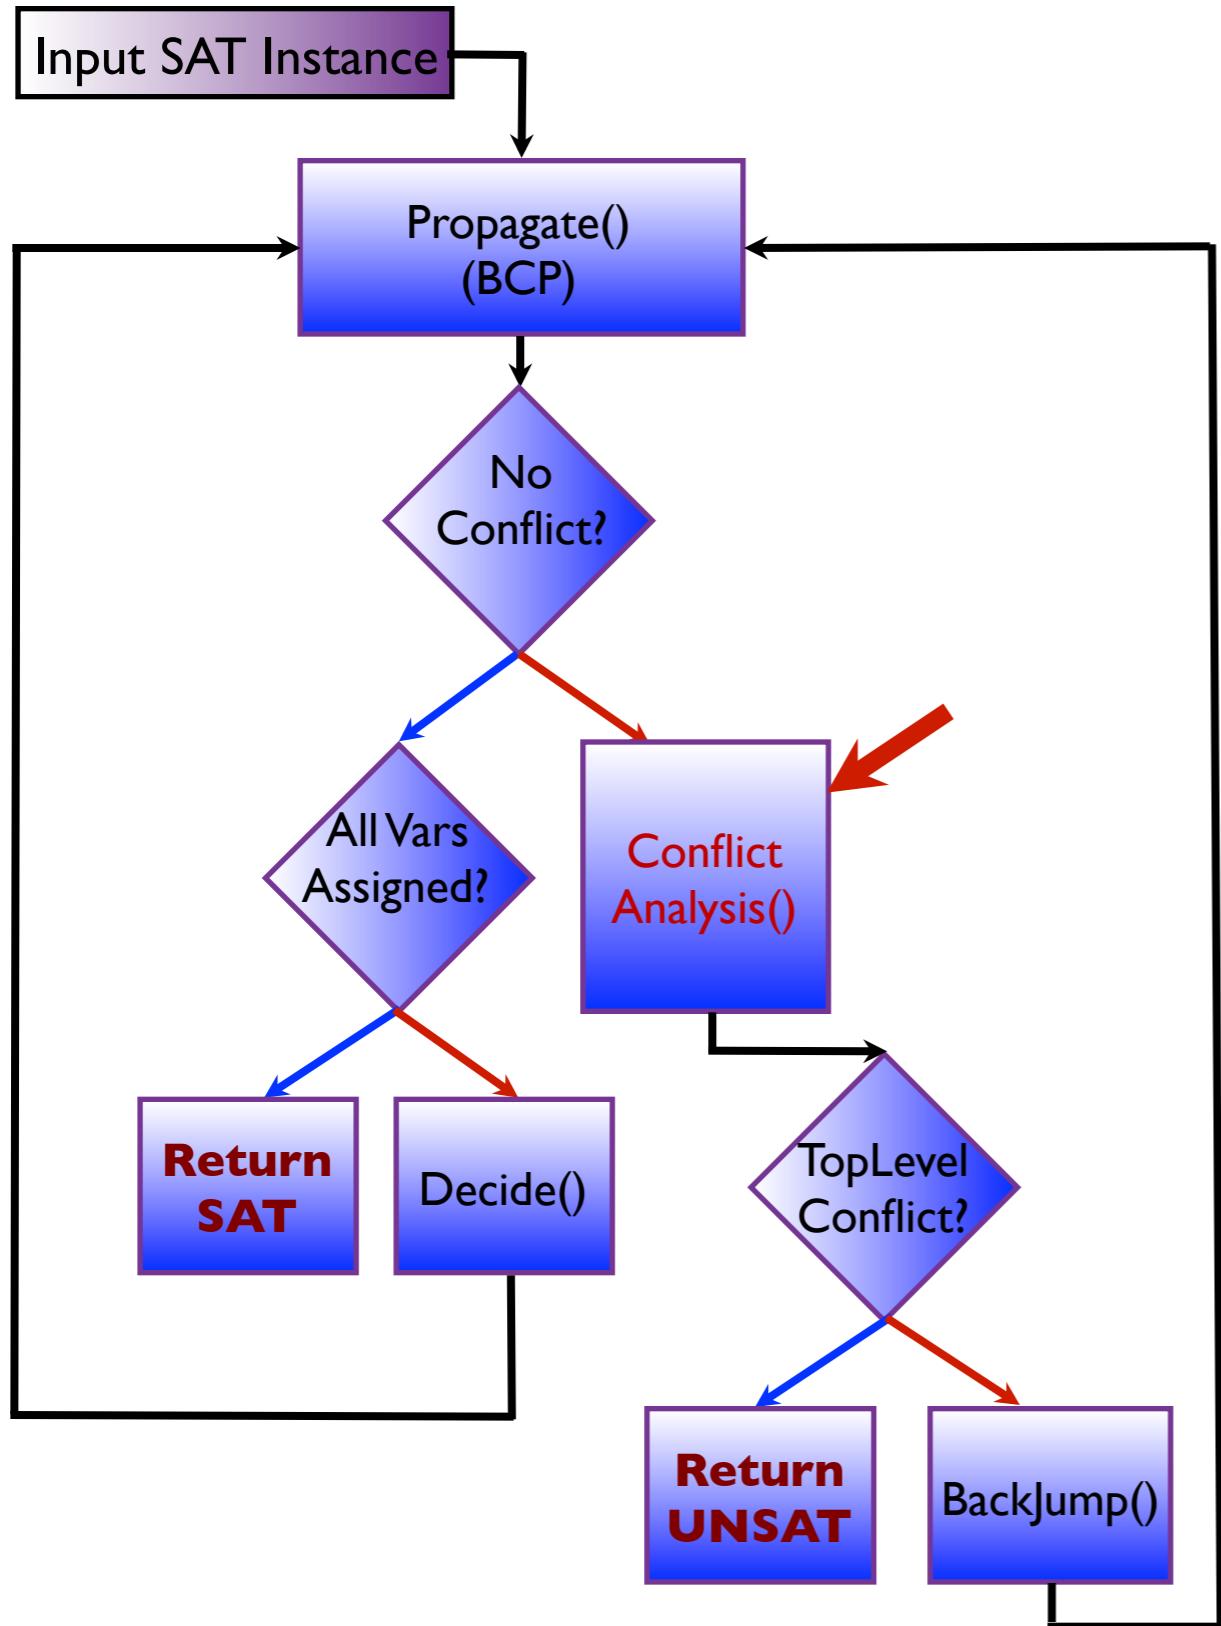
- Conflict: partial assignment is not satisfying

### Decide (Branch)

- Choose a variable & assign some value (**decision**)
- Basic mechanism to do search
- Imposes dynamic variable order
- Least understood aspect of modern SAT solvers

# Modern CDCL SAT Solver Architecture

## Propagate(), Decide(), Analyze/Learn(), BackJump()



### Propagate

- Propagate inferences due to unit clauses
- Most time in solving goes into this

### Detect Conflict?

- Conflict: partial assignment is not satisfying

### Decide (Branch)

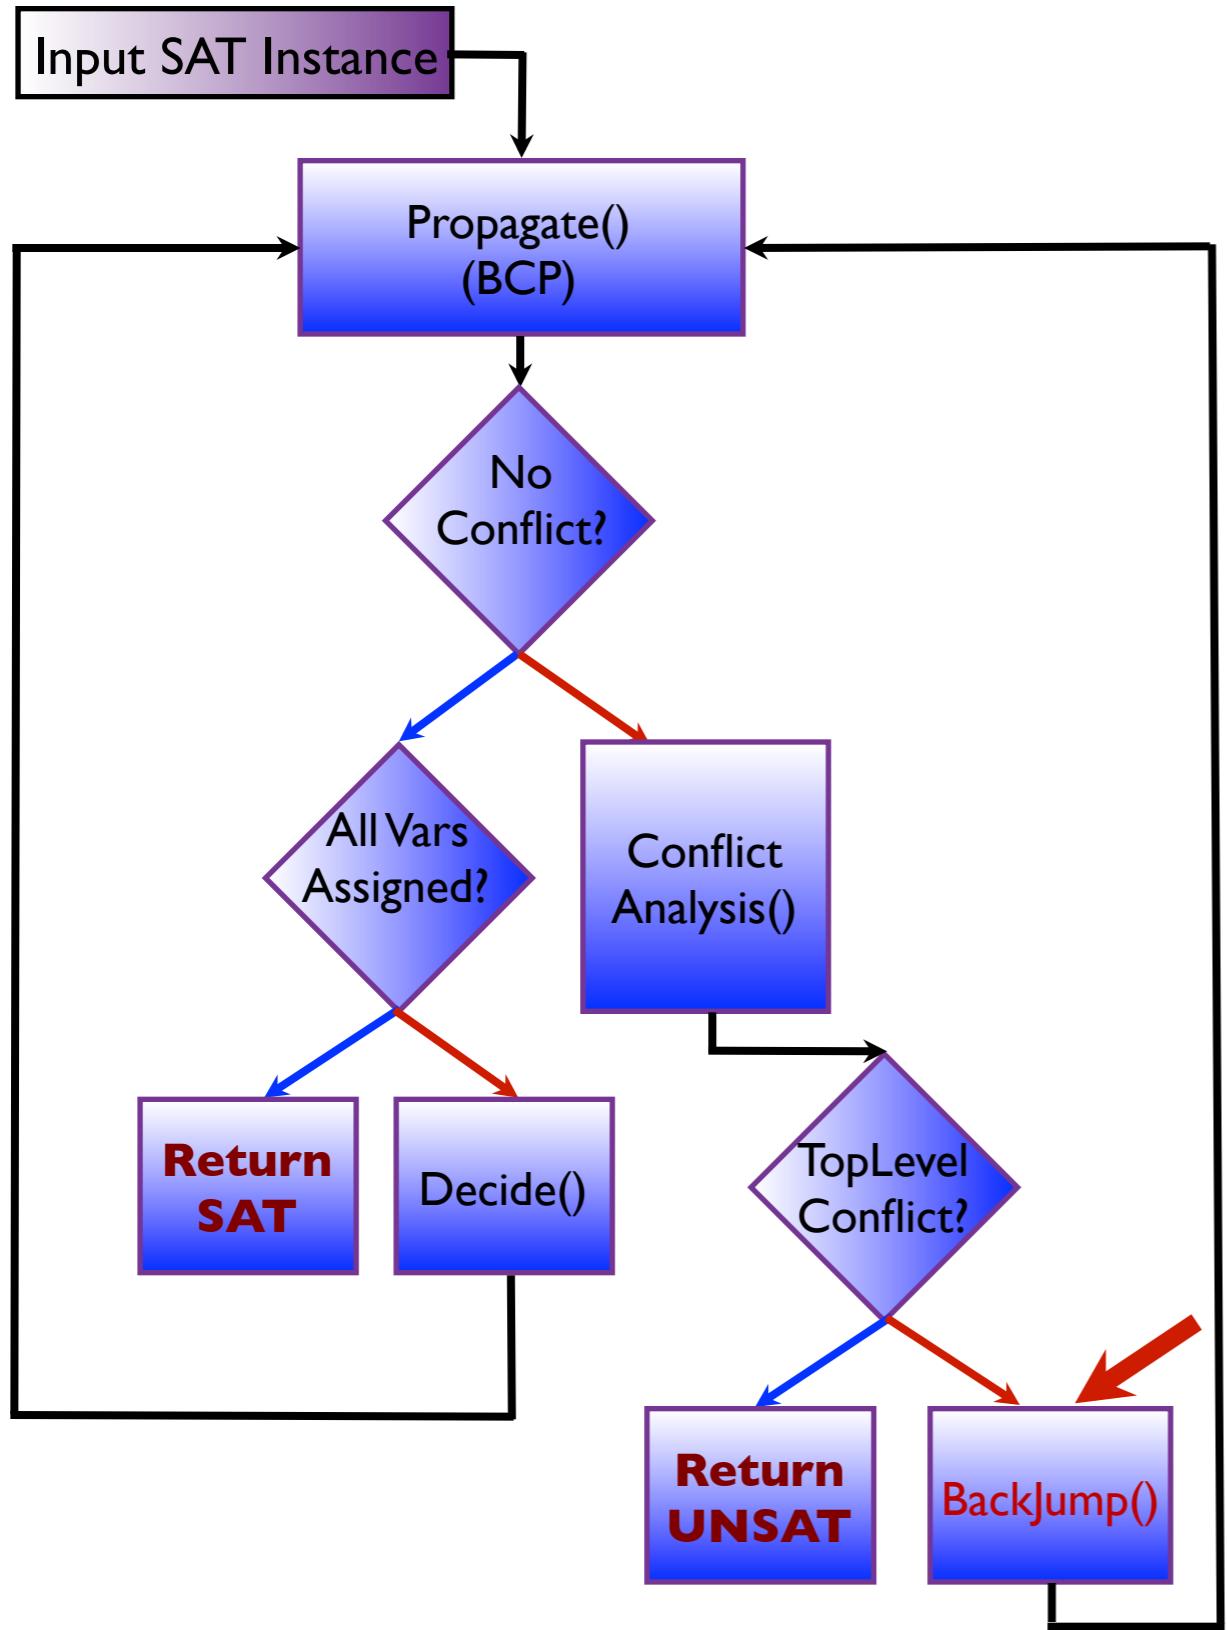
- Choose a variable & assign some value (**decision**)
- Each decision is a decision level
- Imposes dynamic variable order
- Decision Level (**DL**): variable  $\Rightarrow$  natural number

### Conflict analysis and clause learning

- Compute implications that lead to conflict (**analysis**)
- Construct learnt clause that blocks the non-satisfying and a set of other ‘no-good’ assignments (**learning**)
- Originally developed by Marques-Silva & Sakallah(1996)

# Modern CDCL SAT Solver Architecture

## Propagate(), Decide(), Analyze/Learn(), BackJump()



### Propagate

- Propagate inferences due to unit clauses
- Most time in solving goes into this

### Detect Conflict?

- Conflict: partial assignment is not satisfying

### Decide

- Choose a variable & assign some value (decision)
- Each decision is a decision level
- Imposes dynamic variable order
- Decision Level (**DL**): variable  $\Rightarrow$  natural number

### Conflict analysis and clause learning

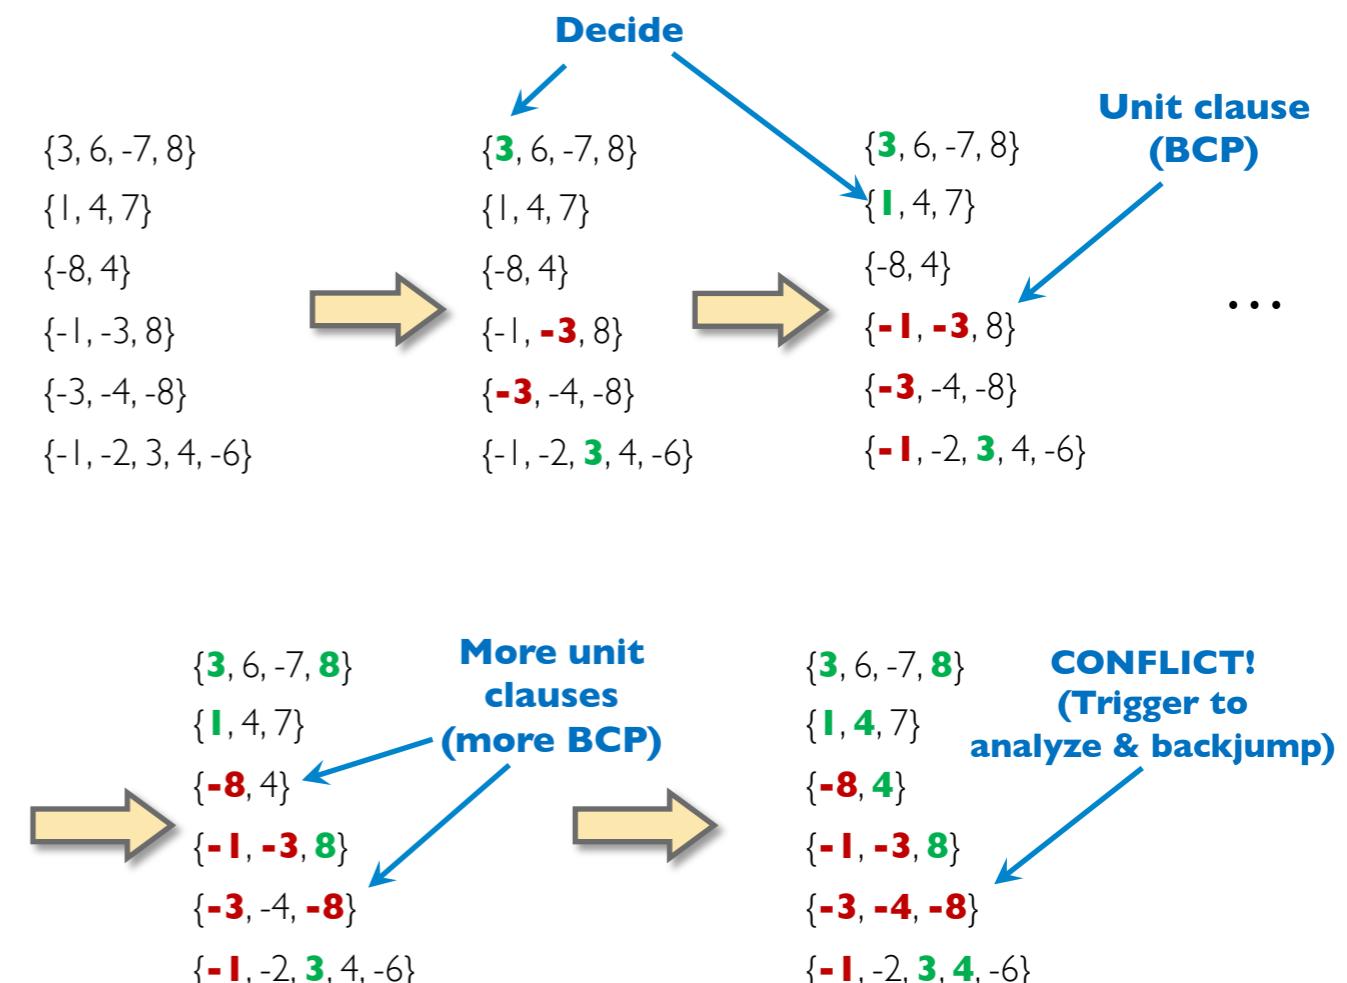
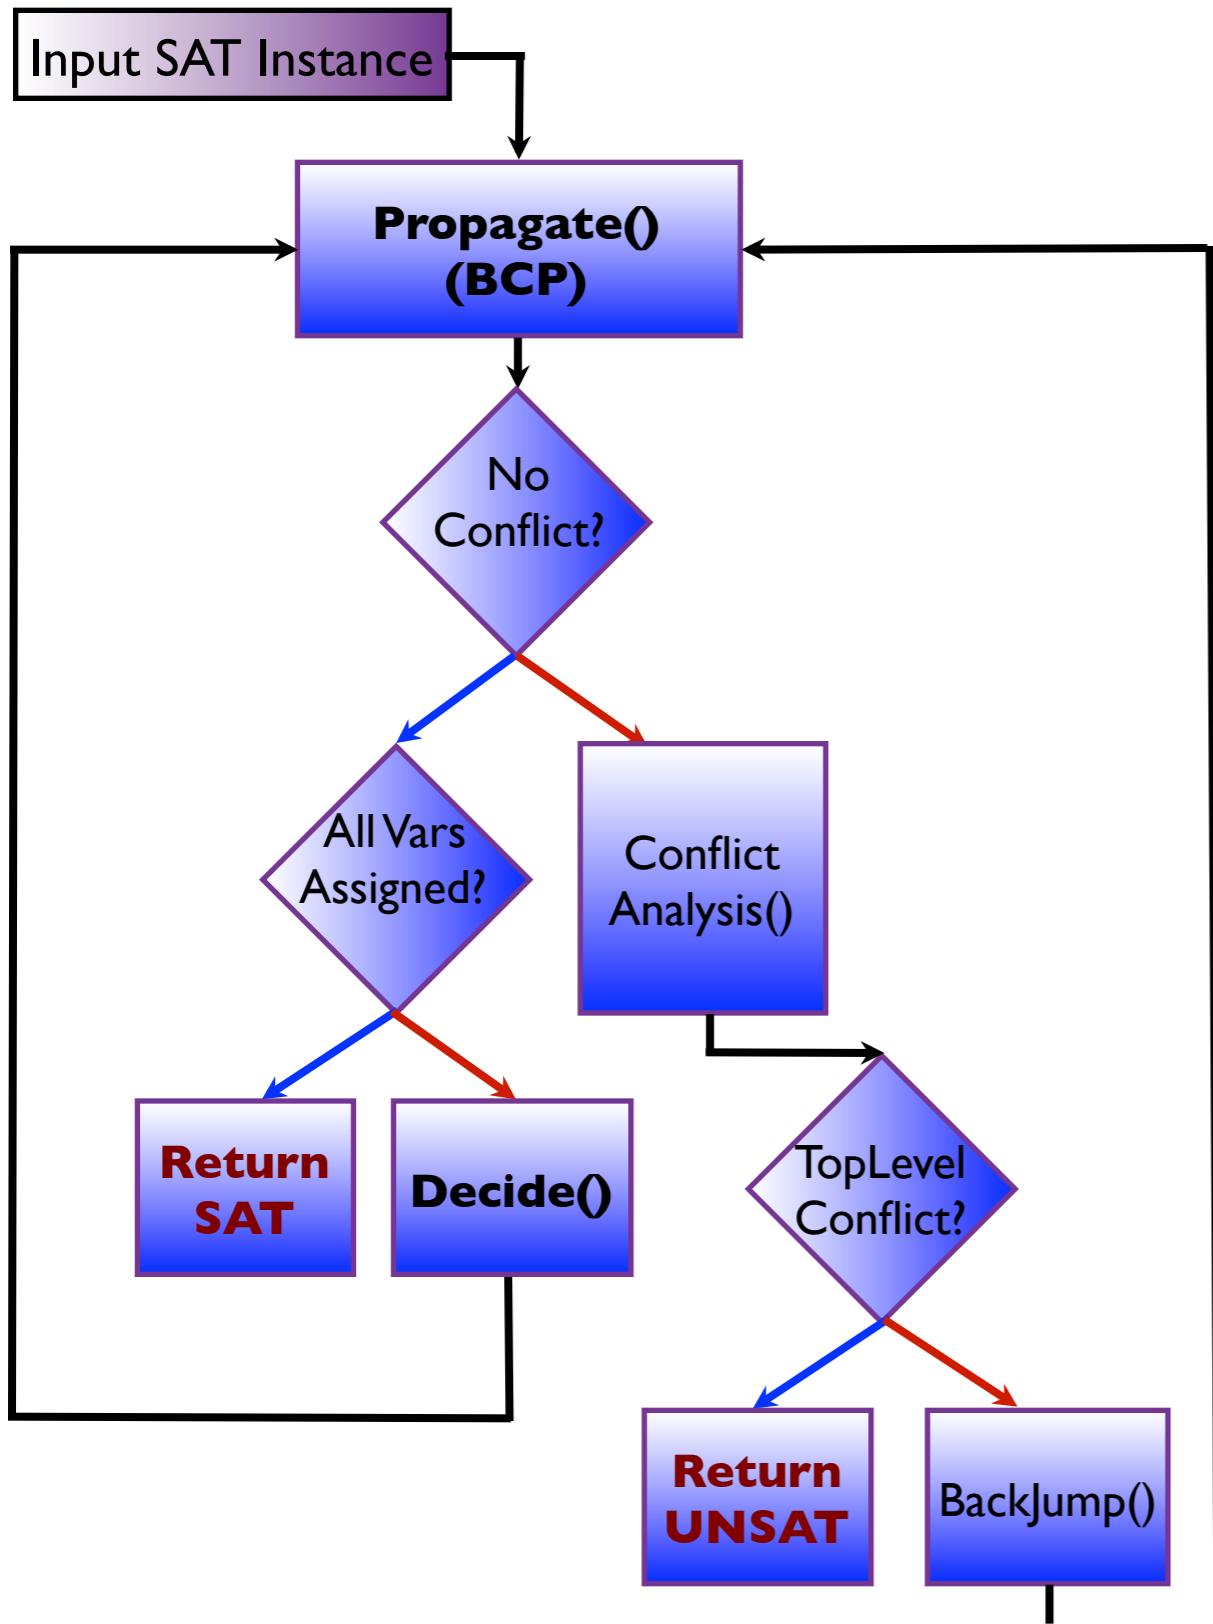
- Compute implications that lead to conflict (**analysis**)
- Construct learnt clause that blocks the non-satisfying and a set of other ‘no-good’ assignments (**learning**)

### Conflict-driven BackJump

- Undo some decision(s) that caused no-good assignment
- Assign undone variable different value
- Go back several decision levels
- The technique is sound, i.e., no solutions are missed
- Backjump: Marques-Silva, Sakallah (1999)
- Backtrack: Davis, Putnam, Loveland, Logemann (1962)

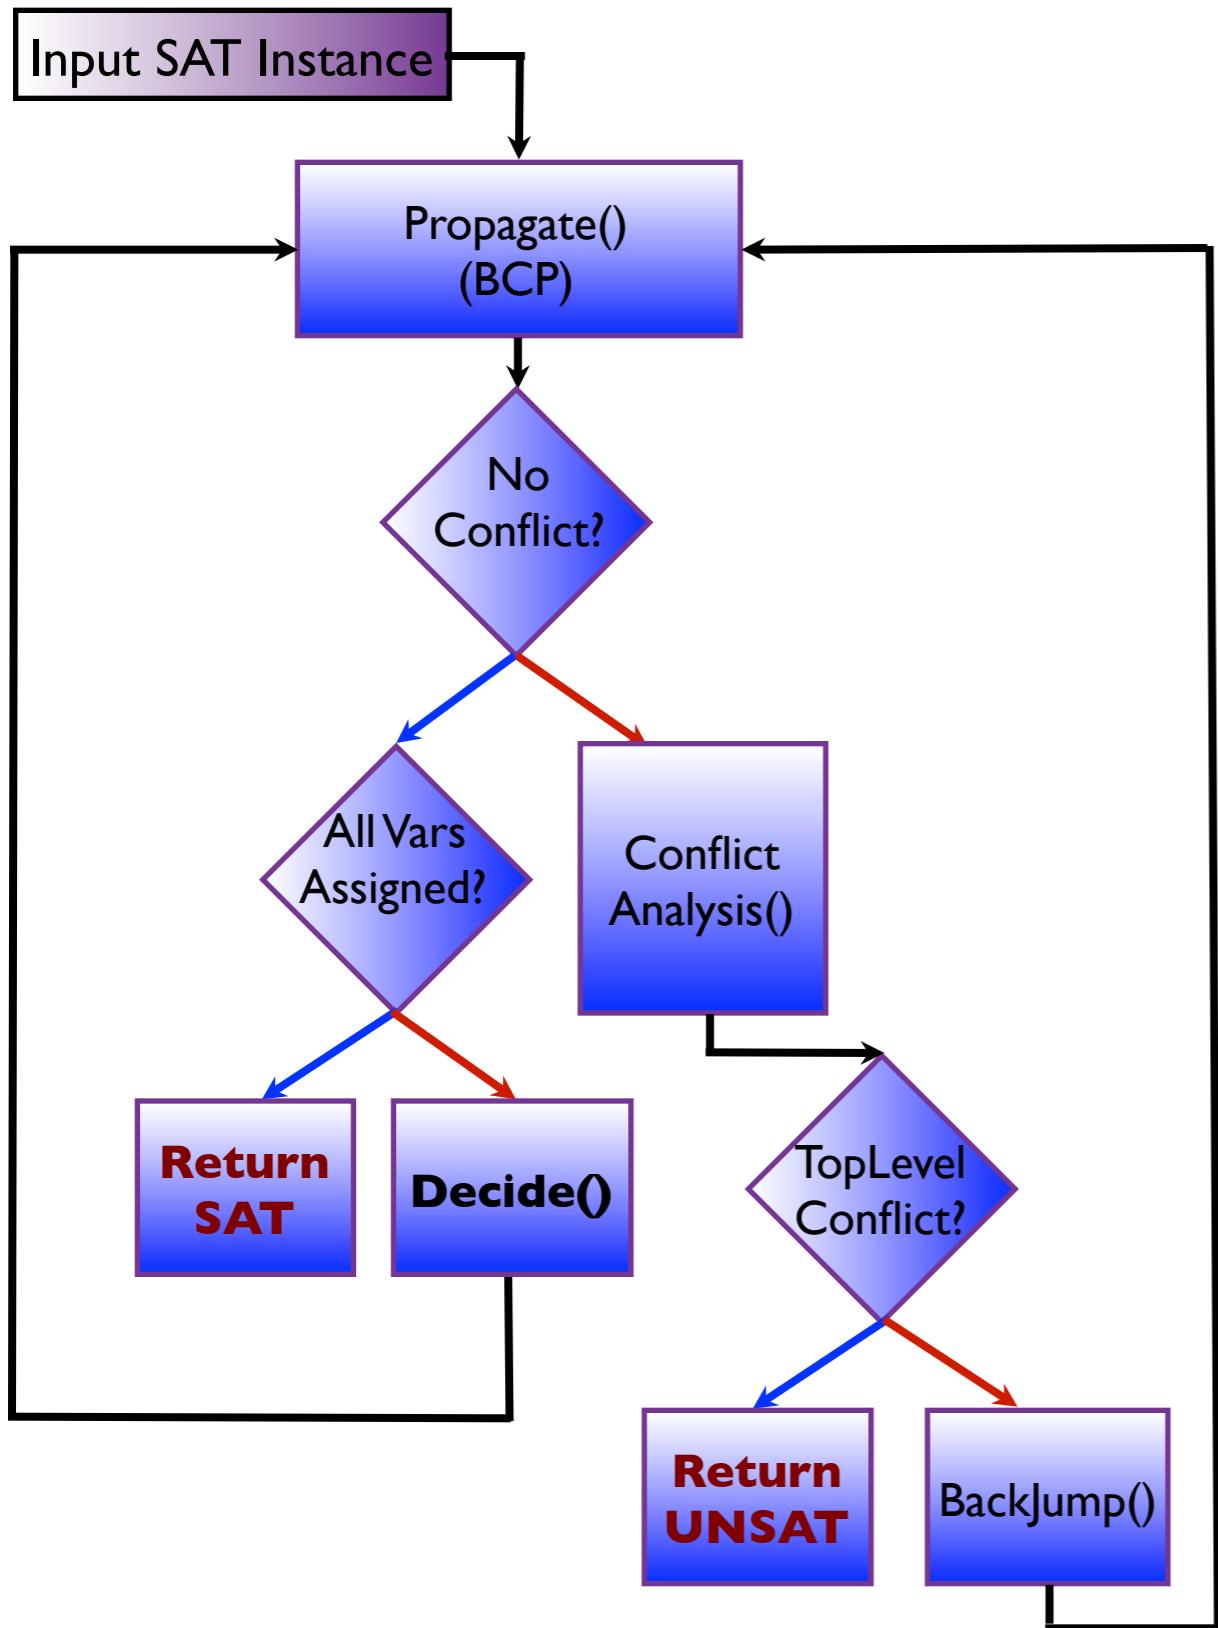
# Modern CDCL SAT Solver Architecture

## Propagate(), Decide(), Analyze/Learn(), BackJump()



# Modern CDCL SAT Solver Architecture

## Decide() Details: VSIDS Heuristic



### Decide() or Branching()

- Choose a variable & assign some value (decision)
- Imposes dynamic variable order

### How to choose a variable

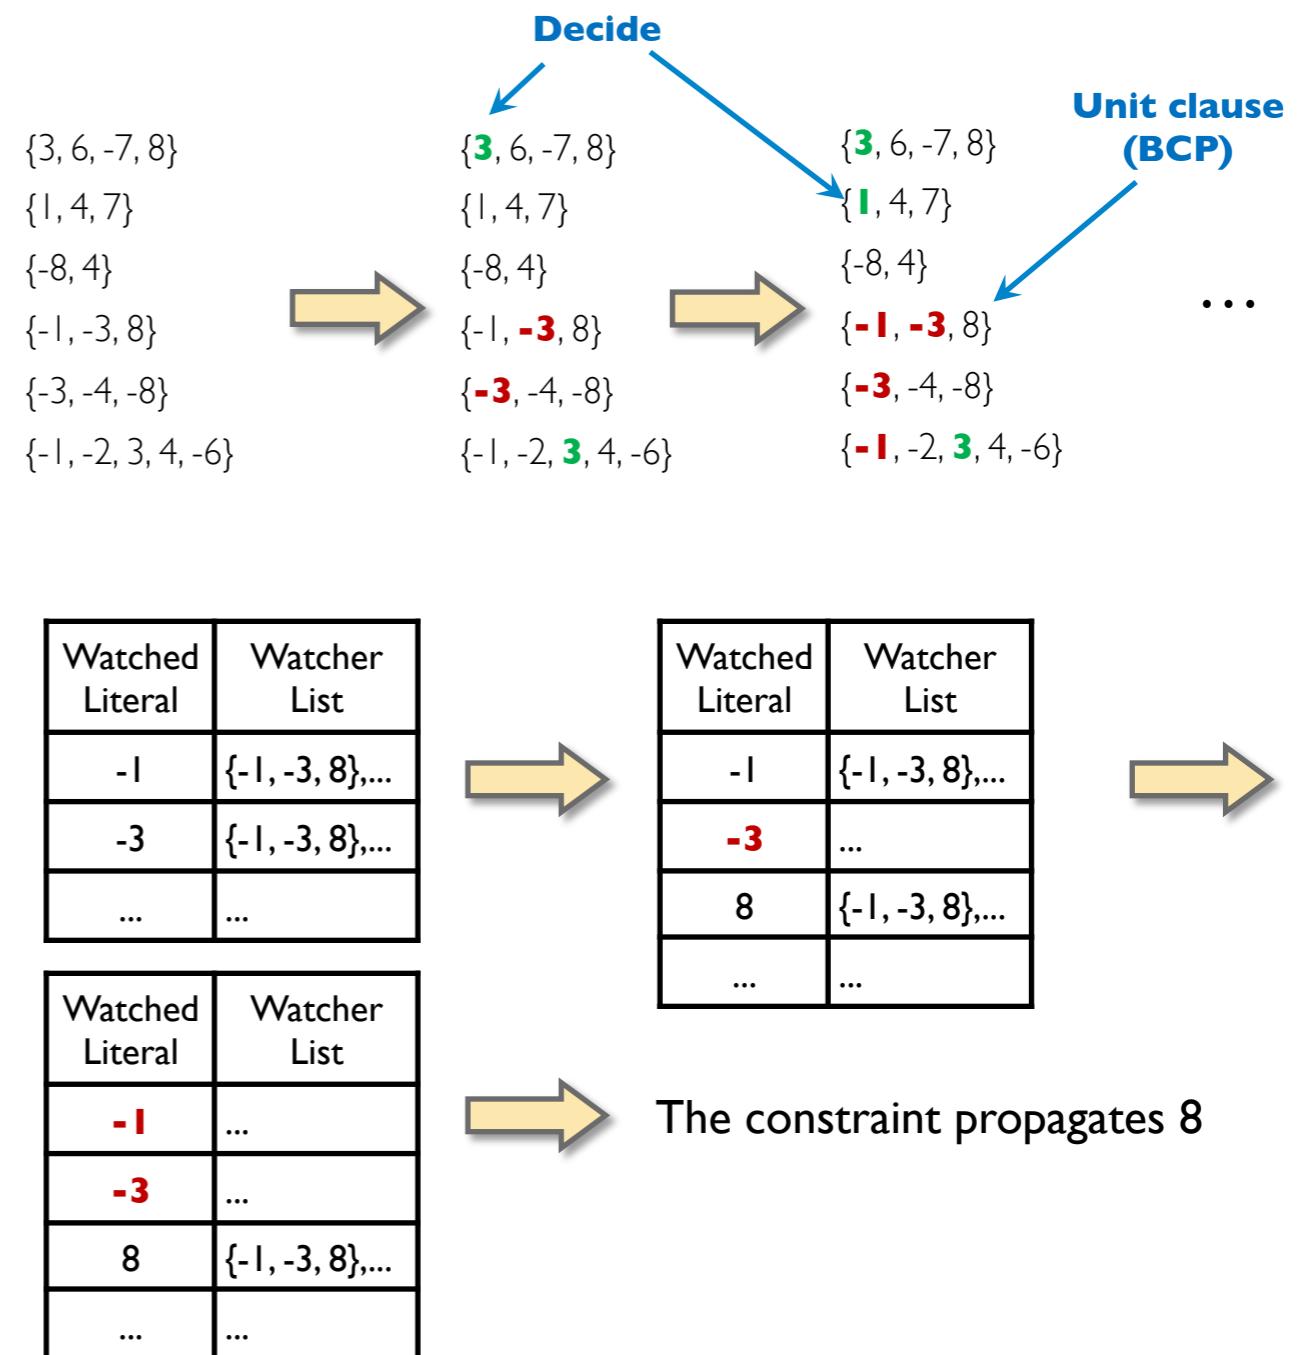
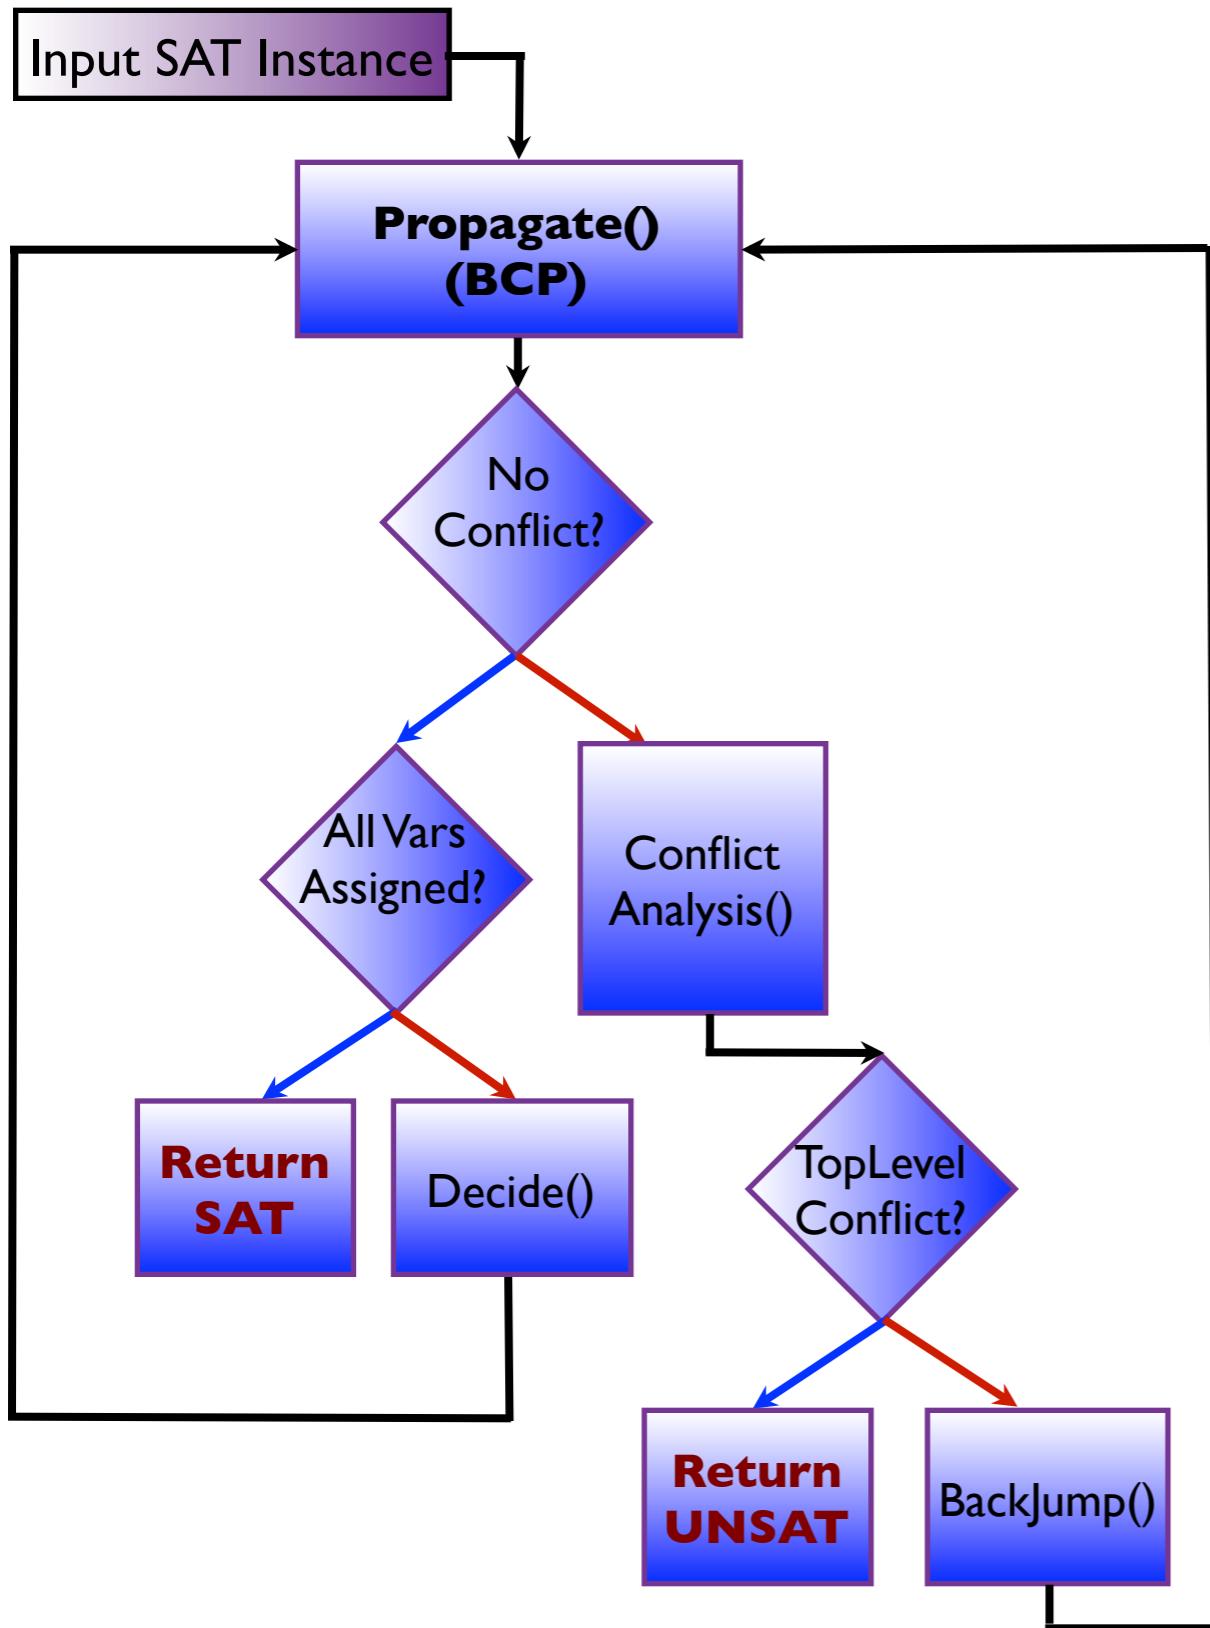
- VSIDS heuristics (zChaff team, DAC 2001)
- Each variable is assigned an **activity value (a real number)**
- Activity is **bumped additively**, if variable occurs in the “most recent” learnt clause
- At regular intervals, the activity of all variables is **decayed** by **multiplying** by a constant less than 1
- The “next decision variable” is the variable with highest activity
- One of the least understood aspects of SAT solving

### Why does VSIDS work so well (ongoing work)

- Imposes dynamic variable order
- Behaves like a low-pass filter or an Exponential Moving Average
- Favors variables that are recently persistently in “conflict”
- More work required to understand the power of VSIDS

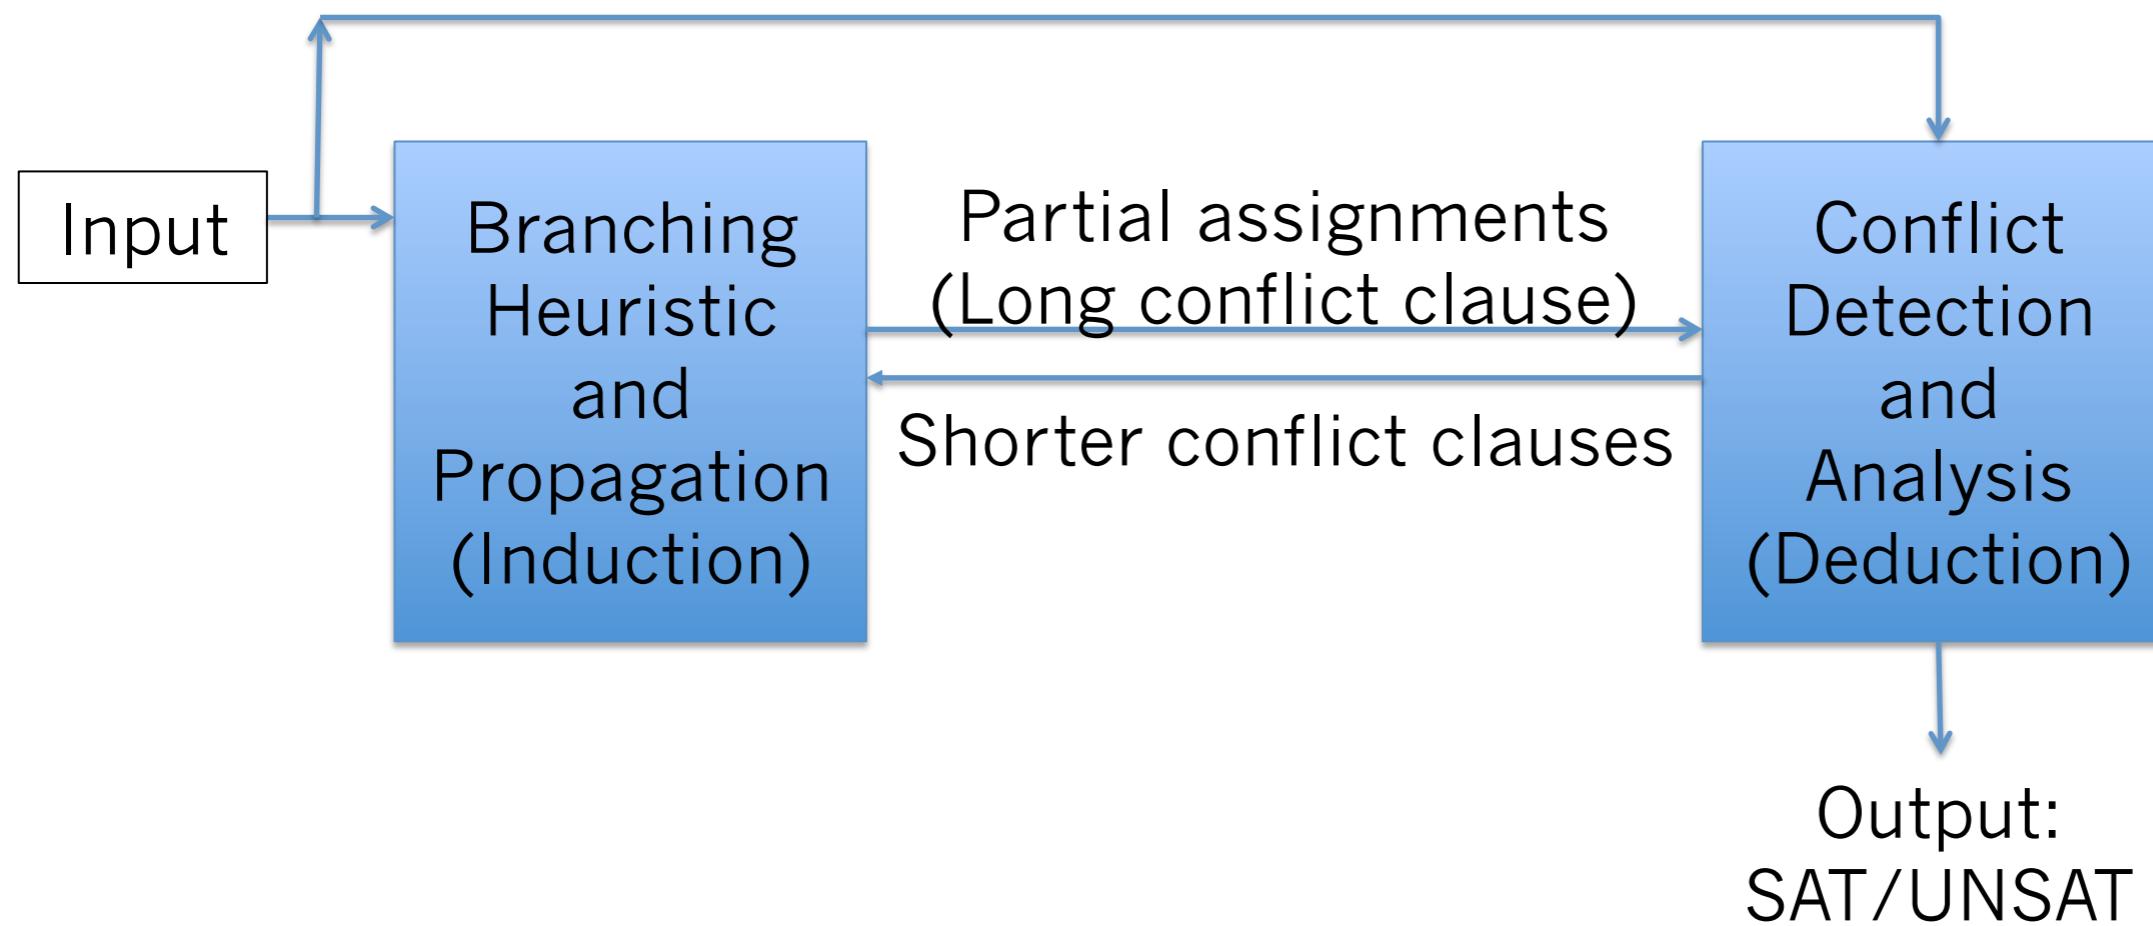
# Modern CDCL SAT Solver Architecture

## Propagate() Details: Two-watched Literal Scheme



# A Model for CDCL Solvers

## CDCL Solvers: Induction and Deduction with Feedback



# Rest of the Talk

## **Topics covered on SAT Solvers**

- Motivation for SAT/SMT solvers in software engineering
- High-level description of the SAT problem
- Key techniques used in SAT solvers
  - DPLL search and its shortcomings
  - Modern CDCL SAT solver: propagate (BCP), decide (VSIDS), conflict analysis, clause learn and backJump, clause deletion
  - Big lesson: learning from mistakes

## **Topics covered on SMT Solvers**

- Modern SMT solvers and key techniques
  - Combination of theories
  - DPLL(T): Boolean reasoning over theories
  - Various forms of abstraction and refinement techniques
- Future directions

## **Impact of Community Structure on SAT Solver Performance (presented at SAT 2014)**

## The SMT Problem

# Satisfiability over Boolean Combination of First-order Theories

- SMT stands for Satisfiability Modulo Theories
- An SMT formula is a Boolean combination of formulas over first-order theories
- Example of SMT theories include bit-vectors, arrays, integer and real arithmetic, strings,...
- The satisfiability problem for these theories is typically hard in general (NP-complete, PSPACE-complete,...)

# Motivations for SMT Solvers

## Aren't SAT Solvers Enough?

- Program semantics are easily expressed over these theories
- Many software engineering problems can be easily reduced to the SAT problem over first-order theories
- Reduction to Boolean SAT not always a good option because of “blowup” during reduction
- SMT expressions have structure that can be more easily detected and solved using powerful heuristics, than by analyzing the equivalent Boolean representation

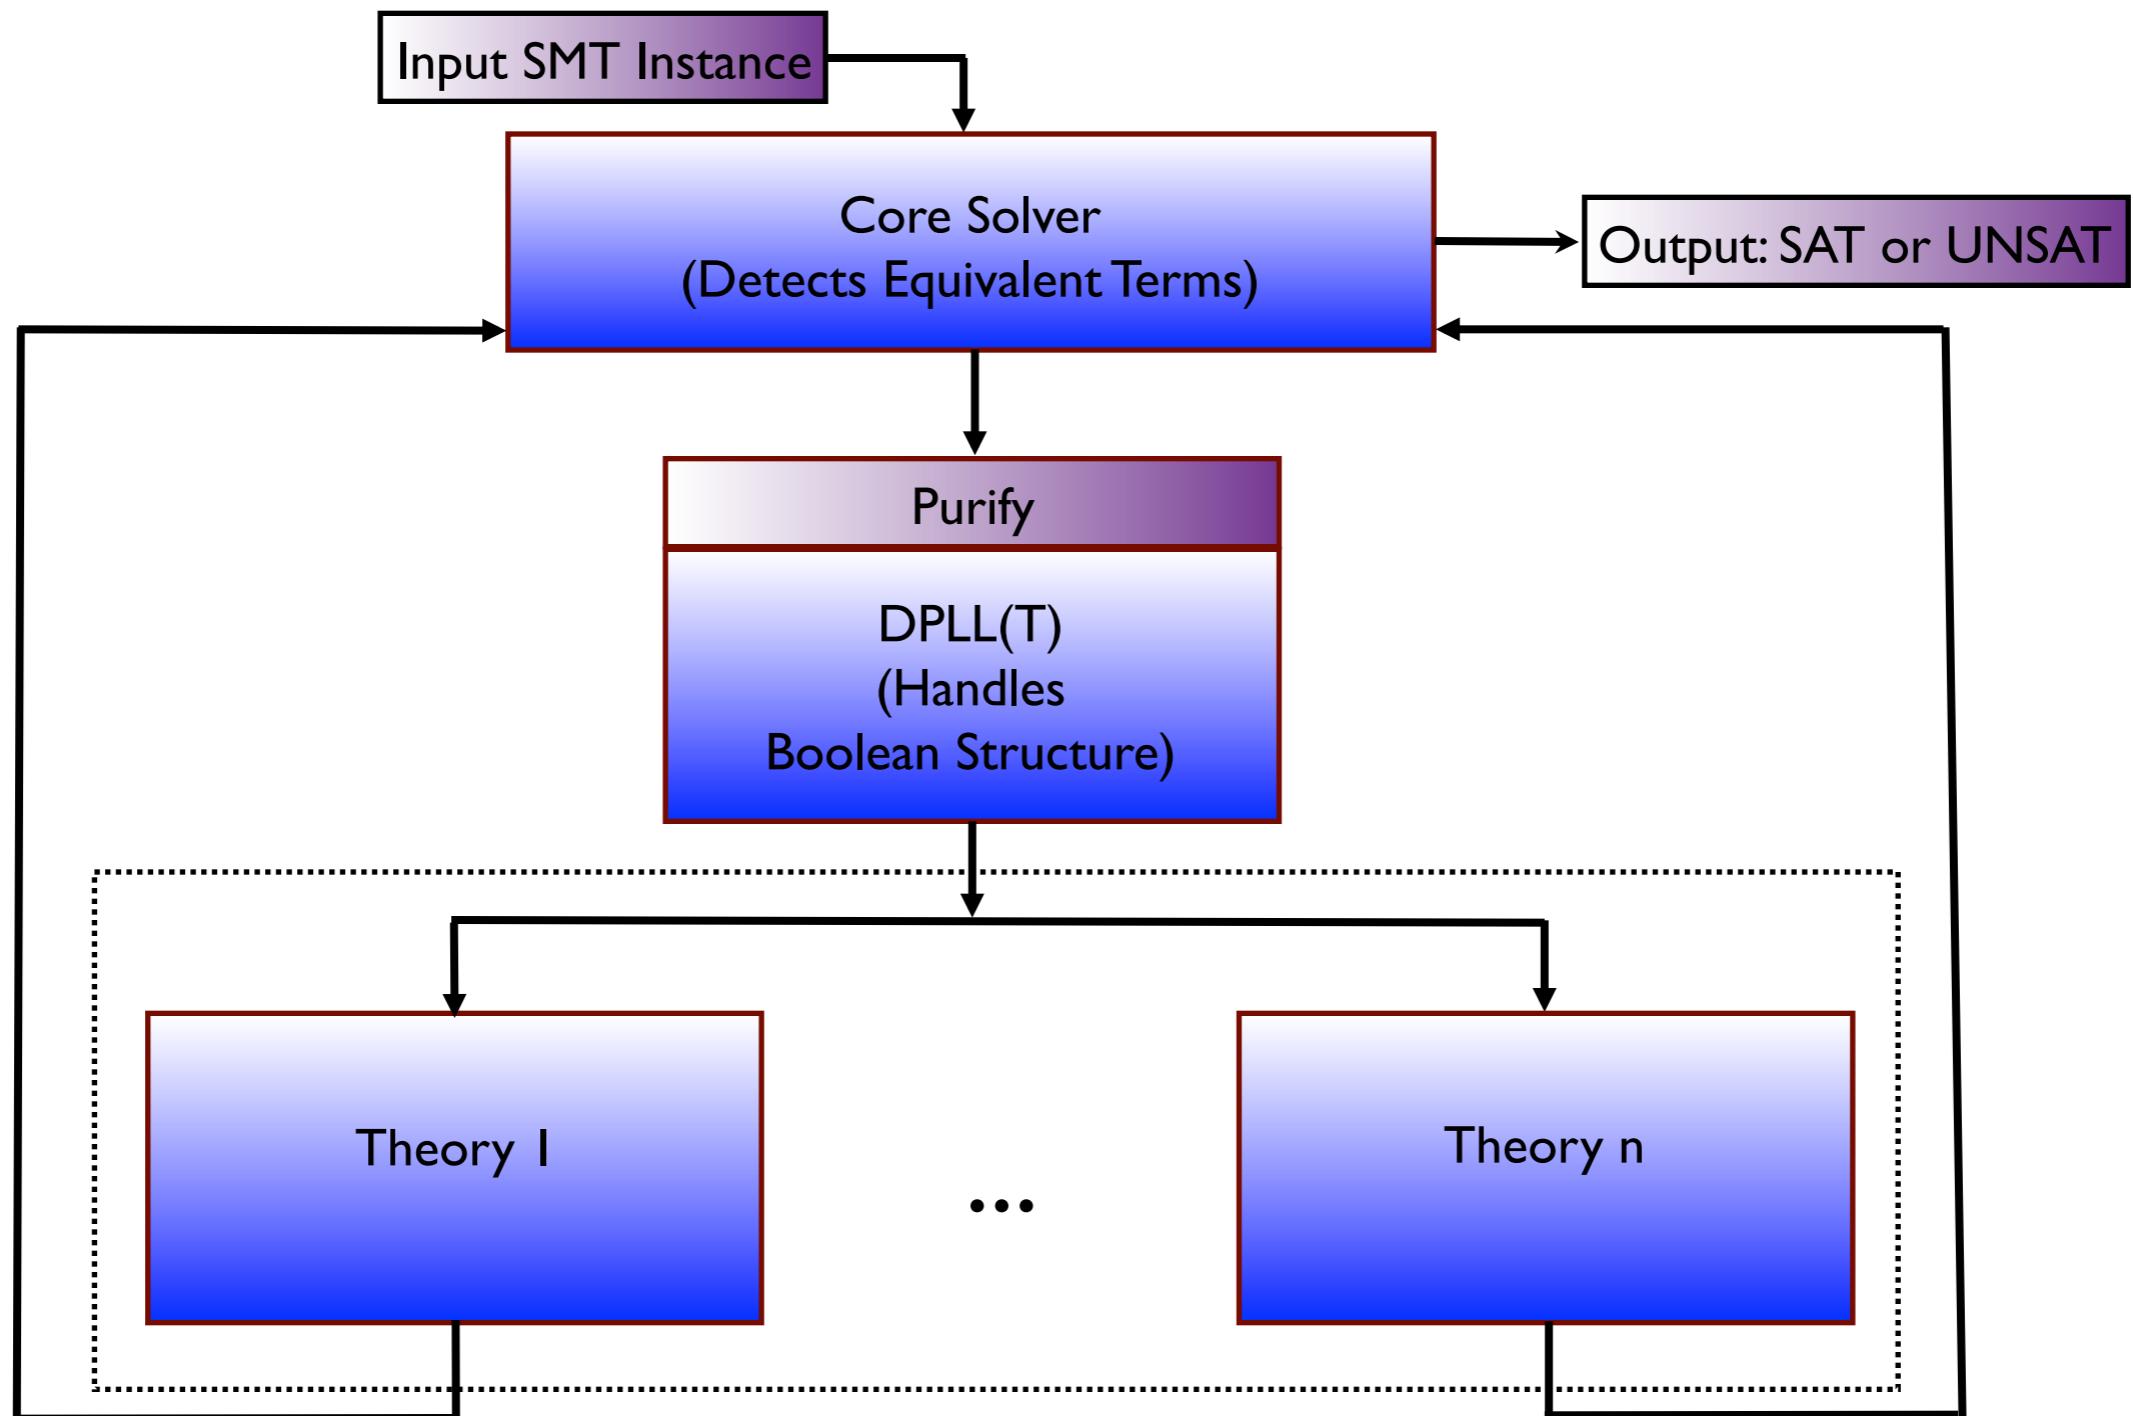
# Modern SMT Solvers

## Key Design Concepts

- Combination of theories (Nelson-Oppen, Shostak, Ghilardi,...)
- DPLL( $T$ ): lazily solve Boolean combination of literals from “different” theories
- Combining “lazy (abstraction and refinement)” and “eager” techniques to solve SAT problem for individual theories

# Standard-issue SMT Solver Architecture

## Combination of theories & DPLL(T)



# Standard-issue SMT Solver Architecture

## Combination of theories: Nelson-Oppen

### Problem Statement

- Combine theory solvers to obtain a solver for a union theory

### Motivation

- Software engineering constraints span many “natural” first-order theories
- Some first-order theories do have “practically” efficient solvers
- Modular composition of such decision procedures is often deemed more efficient than a monolithic solver (Caution: not always true)

### How

- Setup communication between individual theory solvers
- Communication over shared predicates (equality in the case of Nelson-Oppen)
- Soundness, completeness and termination

# Standard-issue SMT Solver Architecture

## DPLL(T)

### Problem Statement

- Efficiently handle the Boolean structure of the input formula

### Basic Idea

- Use a SAT solver for the Boolean structure & check assignment consistency against a T-solver
- T-solver only supports conjunction of T-literals

### Improvements (Nuiwenheus, Oliveras, and Tinelli 2006)

- Check partial assignments against T-solver
- Do **theory propagation** (similar to SAT solvers)
- **Conflict analysis guided by T-solver** & generate conflict clauses (similar to SAT solvers)
- BackJump (similar to SAT solvers)

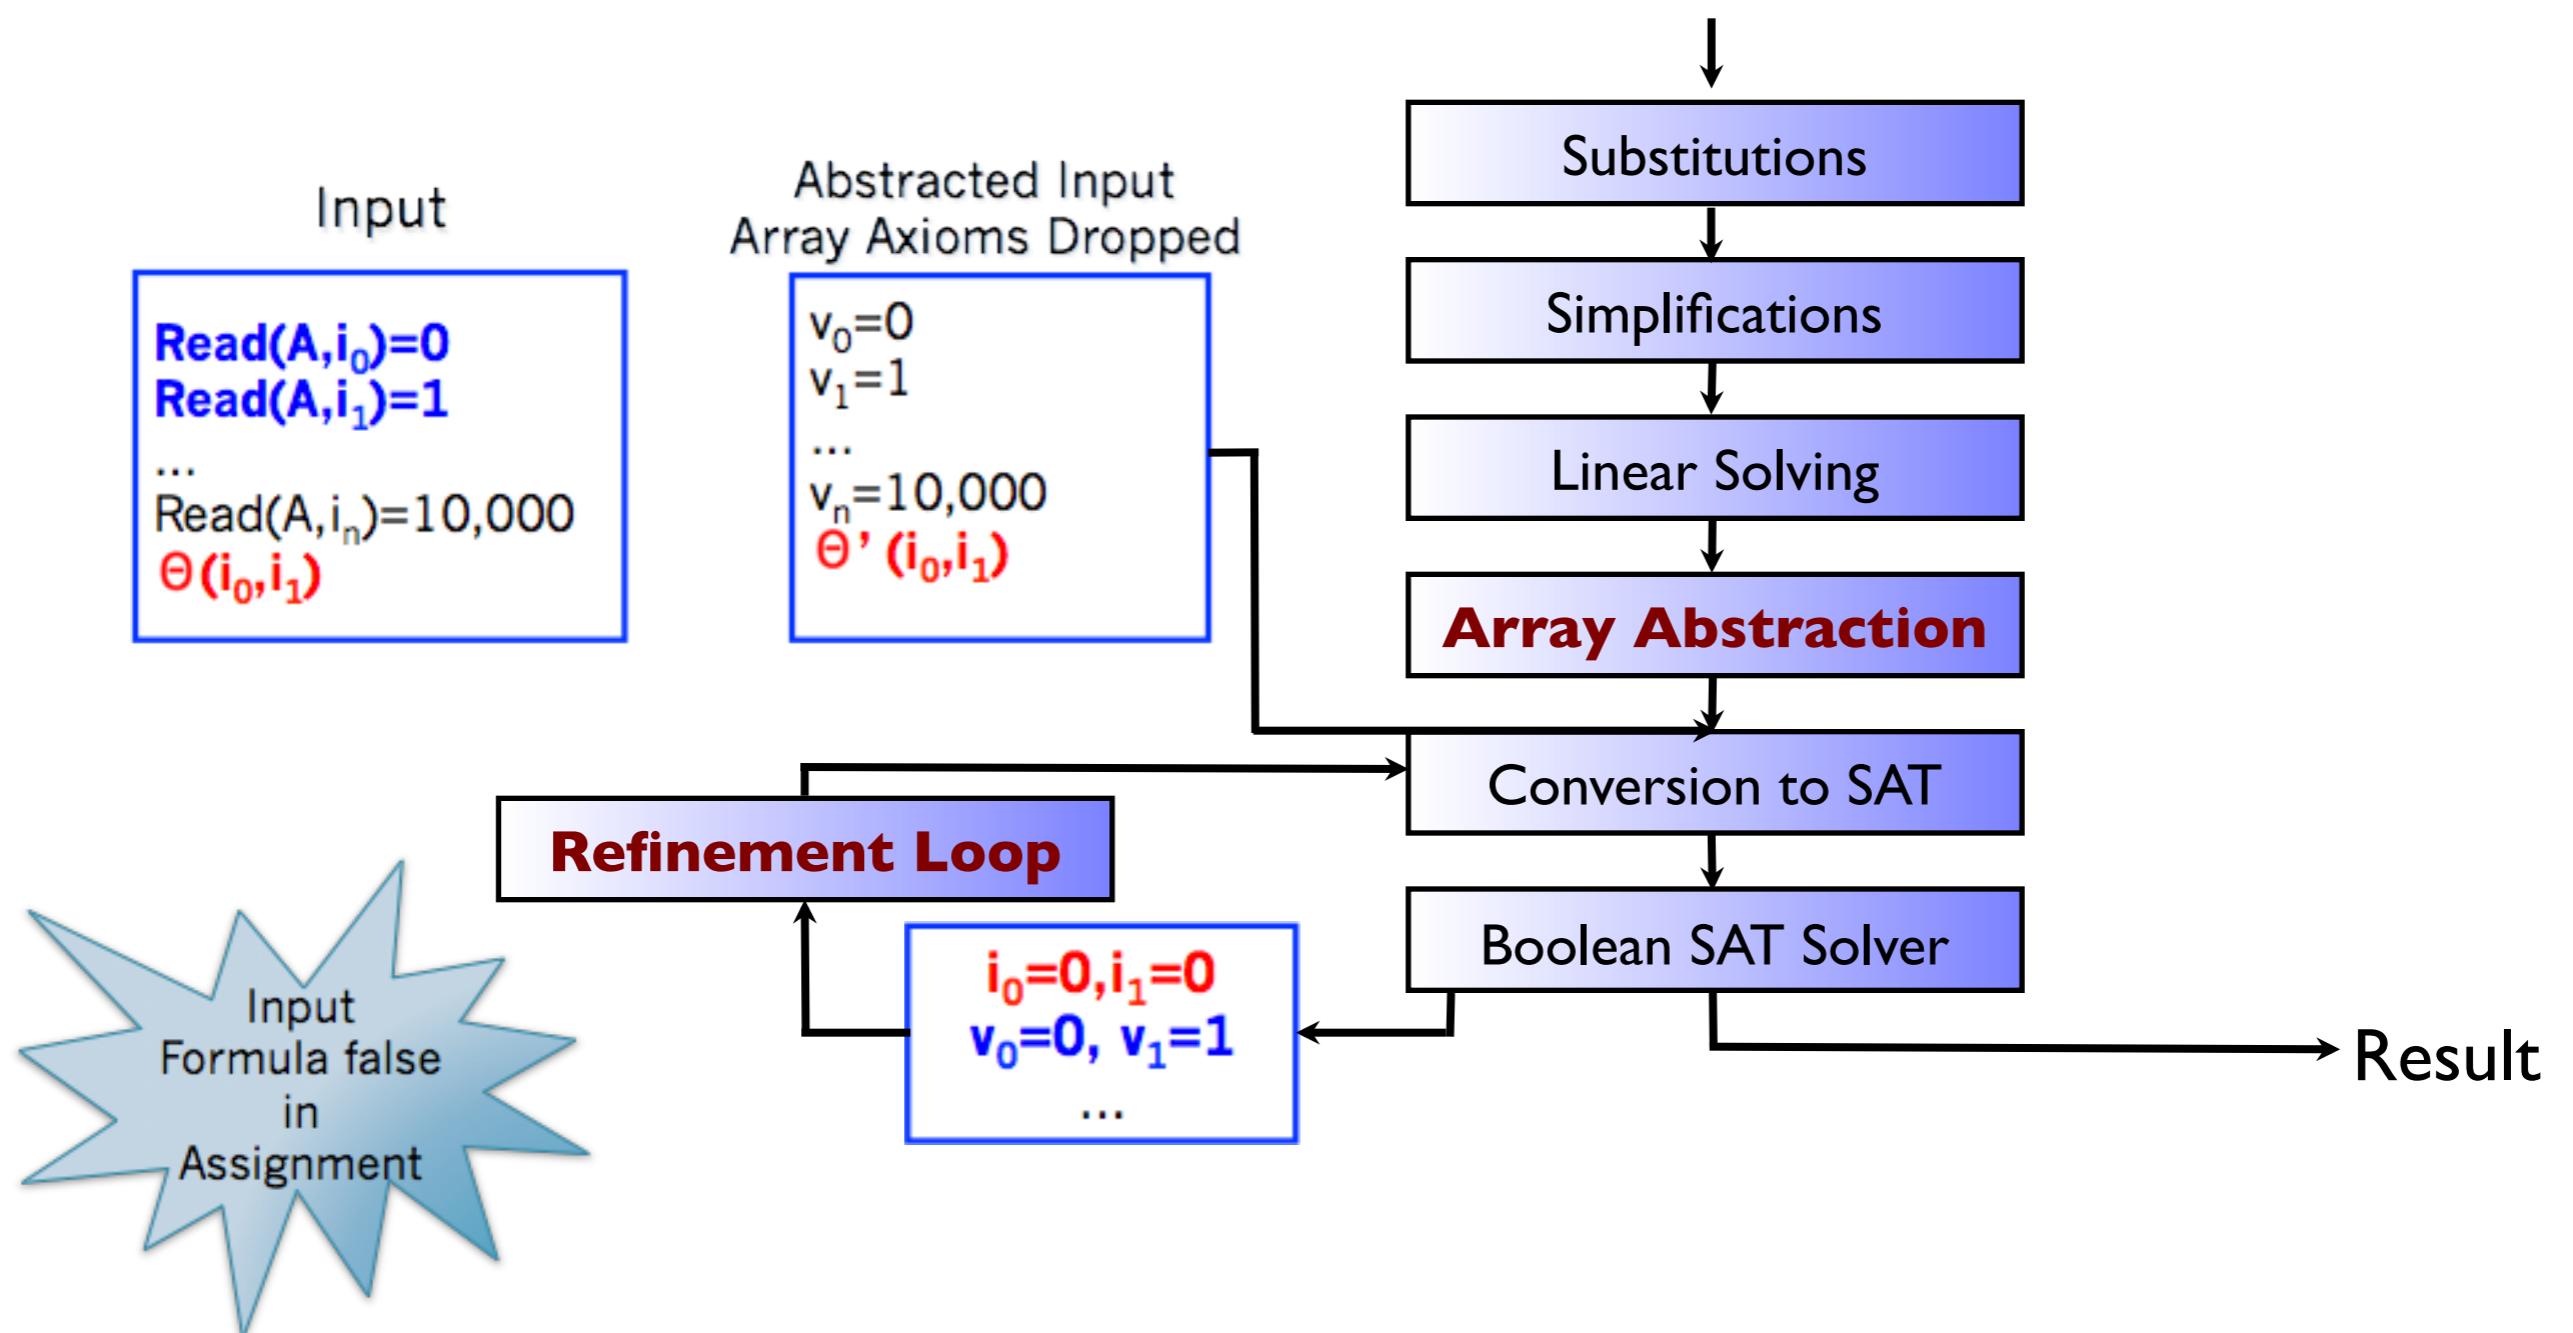
# Modern SMT Solvers

## Key Design Concepts

- Combination of theories (Nelson-Oppen, Shostak, Ghilardi,...)
- DPLL( $T$ ): lazily solve Boolean combination of literals from “different” theories
- Combining “lazy (abstraction and refinement)” and “eager” techniques to solve SAT problem for individual theories

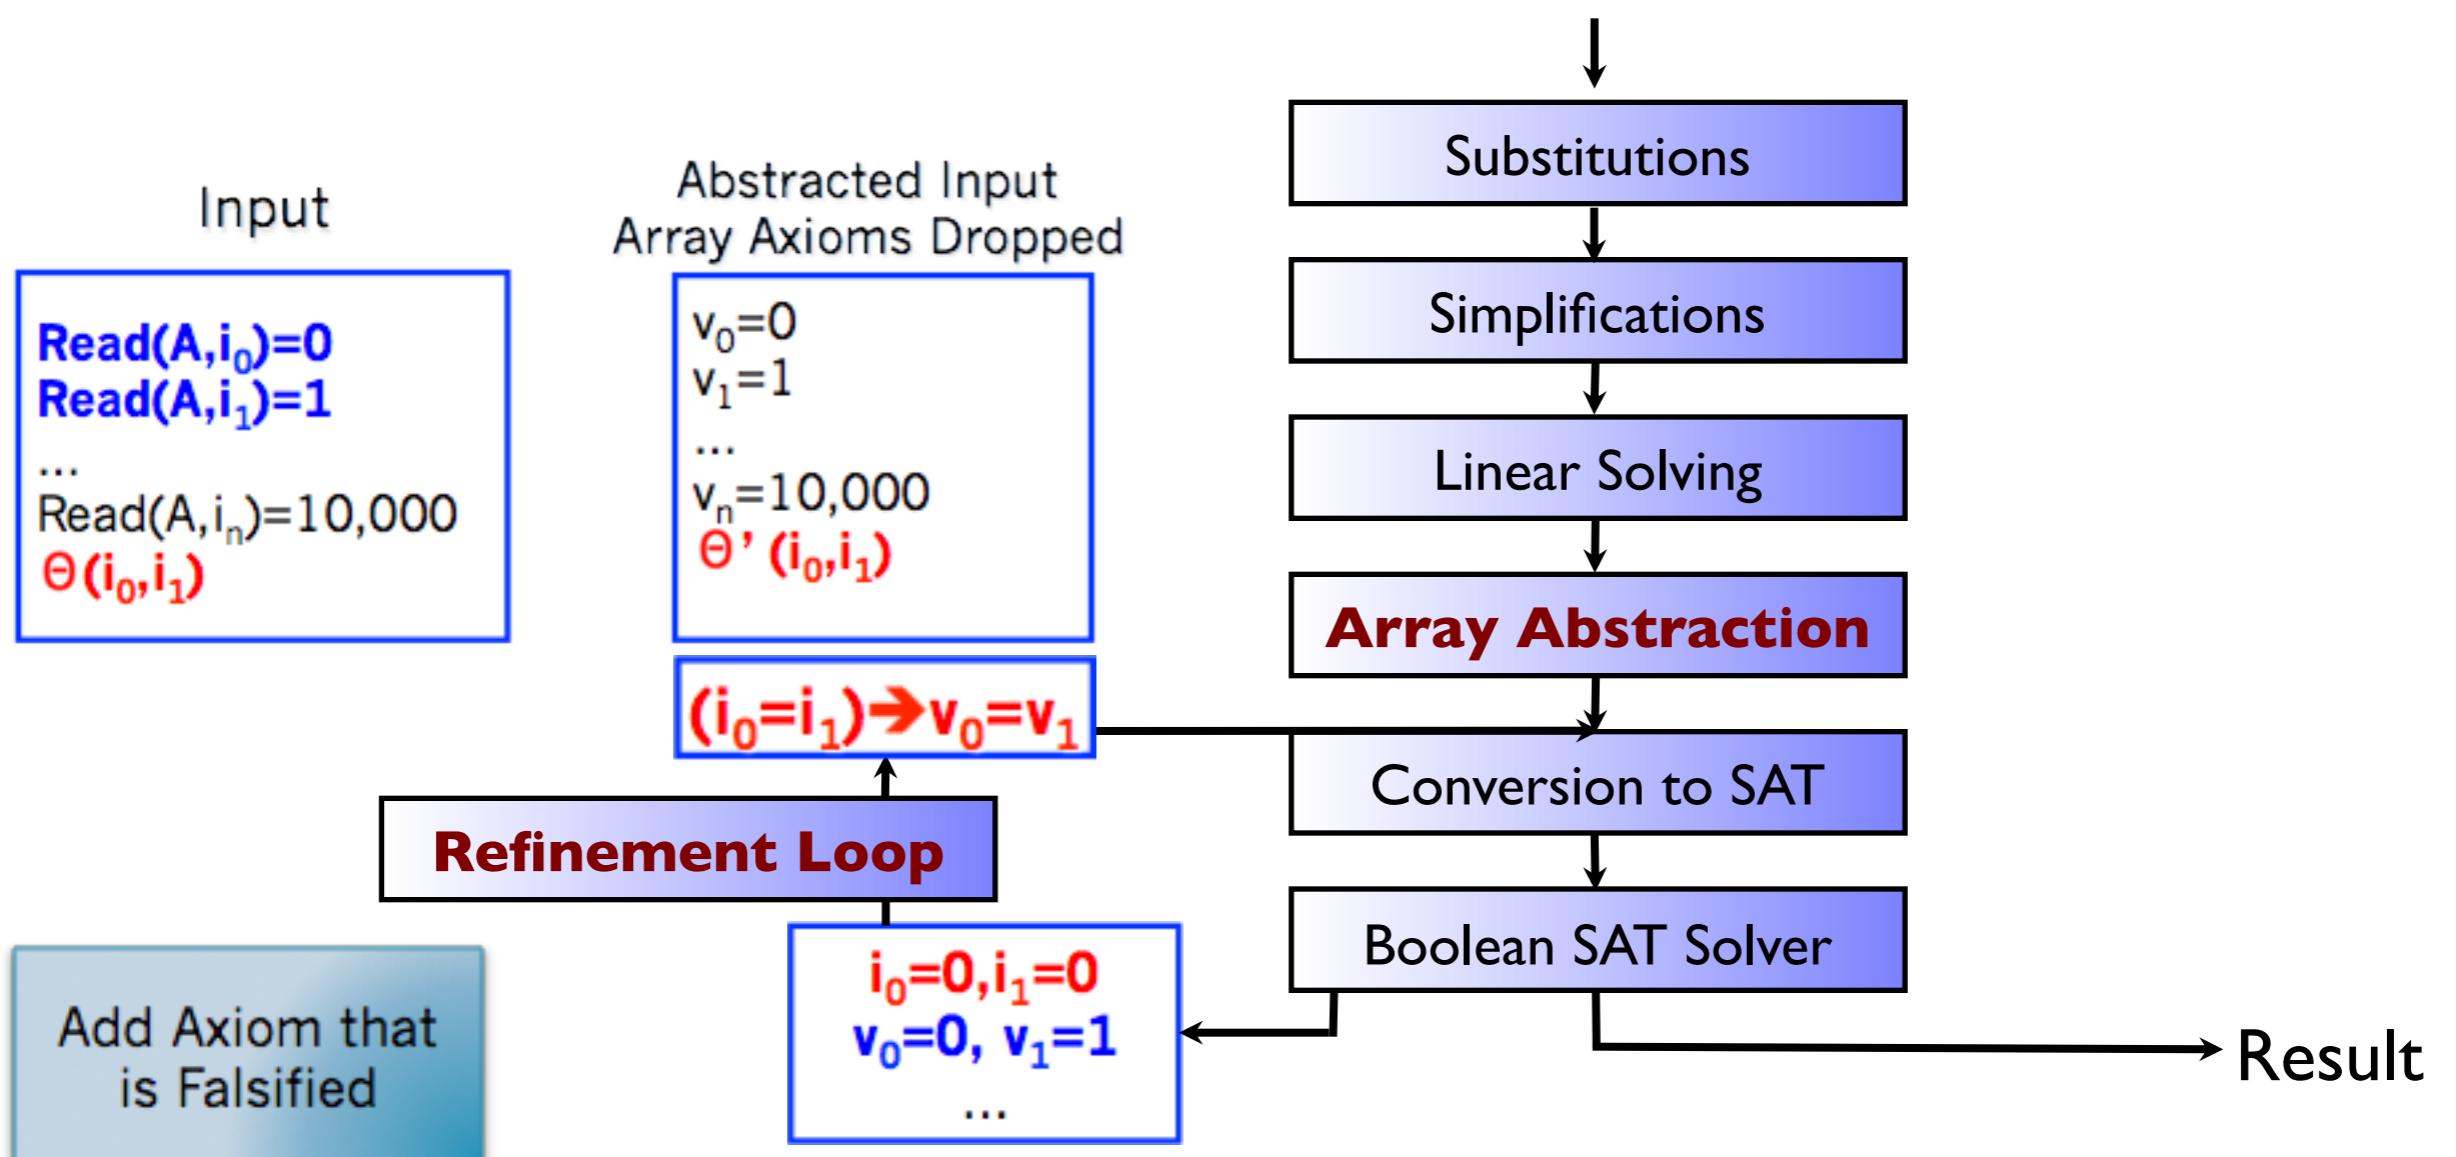
# How STP Works

## Eager for Bit-vectors, Lazy for Arrays



# How STP Works

## Abstraction-refinement for Array-reads



# An Incomplete History of SAT and SMT Solvers

<b>Category</b>	<b>Research Project</b>	<b>Researcher/Institution/Time Period</b>
Theorem Proving (very early roots of decision procedures)	NuPRL Boyer-Moore Theorem Prover ACL2 PVS Proof Checker	Robert Constable / Cornell / 1970's-present Boyer & Moore / UT Austin / 1970's-present Moore, Kauffmann et al. / UT Austin / 1980's - present Natarajan Shankar / SRI International / 1990's-present
SAT Solvers	DPLL GRASP (Clause learning and backjumping) Chaff & zChaff MiniSAT	Davis, Putnam, Logemann & Loveland / 1962 Marques-Silva & Sakallah / U. Michigan / 1996-2000 Zhang, Malik et al. / Princeton / 1997-2002 Een & Sorensson / 2005 - present
Combinations	Simplify Shostak ICS SVC, CVC, CVC-Lite, CVC3 ... Non-disjoint theories	Nelson & Oppen / DEC and Compaq / late 1980s Shostak / SRI International / late 1980's Ruess & Shankar / SRI International / late 1990's Barrett & Dill / Stanford U. / late 1990's Tinelli, Ghilardi, G.,... / 2000 - 2008
DPLL(T)	Barcelogic and Tinelli group	Oliveras, Nieuwenhuis & Tinelli / UPC and Iowa / 2006
Under/Over Approximations	UCLID STP	Seshia & Bryant / CMU / 2004 - present Ganesh & Dill / Stanford / 2005 - present
Widely-used SMT Solvers	Z3 CVC4 OpenSMT Yices MathSAT STP UCLID	DeMoura & Bjorner / Microsoft / 2006 - present Barrett & Tinelli / NYU and Iowa / early 2000's - present Bruttomesso / USI Lugano / 2008 - present Deuterre / SRI International / 2005 - present Cimatti et al. / Trento / 2005 - present Ganesh / Stanford & MIT / 2005 - present Seshia / CMU & Berkeley / 2004 - present

# Rest of the Talk

## Topics covered on SAT Solvers

- Motivation for SAT/SMT solvers in software engineering
- High-level description of the SAT problem
- Key techniques used in SAT solvers
  - DPLL search and its shortcomings
  - Modern CDCL SAT solver: propagate (BCP), decide (VSIDS), conflict analysis, clause learn and backJump, clause deletion
  - Big lesson: learning from mistakes

## Topics covered on SMT Solvers

- Modern SMT solvers and key techniques
  - Combination of theories
  - DPLL(T): Boolean reasoning over theories
  - Various forms of abstraction and refinement techniques
- Future directions

## **Impact of Community Structure on SAT Solver Performance (presented at SAT 2014)**

# Problem Statement

## Why are SAT Solvers efficient for Industrial Instances

- Conflict-driven clause learning (CDCL) Boolean SAT solvers are remarkably efficient for large industrial instances
- This is true for industrial instances from a diverse set of applications
- These instances may have tens of millions of variables and clauses

# Motivation

## Why are SAT Solvers efficient for Industrial Instances

- This question has stumped both theoreticians and practitioners alike
- May lead to better solvers and predictors
- In turn may lead to better solver applications

# Previous Work

## Why are SAT Solvers efficient for Industrial Instances

- Results on the practical front to-date
  - Phase-change at clause-var ratio of 4.25 for randomly-generated instances [GW00]
  - Empirically observed \*only\* for randomly-generated instances
  - Doesn't explain the power of CDCL solvers for industrial instances
  - Industrial SAT instances community structure by Ansotegui, Levy et al.[ALI12, ALI13]
  - Large successful predictive model by Xu, Hoos et al. Basis for a machine learning based predictor [XHHL08]

# Structure in Industrial Instances

## Why are SAT Solvers efficient for Industrial Instances

- SAT solvers exploit structure Inherent to SAT instances
- Relevant questions
  1. What structure of industrial instances do solvers exploit?
  2. What evidence connects so-called structure and solver performance?
  3. How does the solver go about exploiting this structure?
- In this talk we discuss answers to Questions 1 and 2

# Community Structure and SAT Solver Performance

## Take Home Message

- Take-home Message
  - Community structure (the quality of which is measured using a metric called  $Q$ ) of SAT instances strongly affect solver performance
- Working towards
  - Small predictive set of features that forms the basis for a complete explanation (we don't have it yet)

# Community Structure and SAT Solver Performance

---

## Empirical Results

- **Result #1:** Hard random instances have low Q ( $0.05 \leq Q \leq 0.13$ )
- **Result #2:** Number of communities and Q of SAT instances are more predictive of CDCL solver performance than other measures we considered
- **Result #3:** Strong correlation between community structure and LBD (Literal Block Distance) in Glucose solver

# Community Structure and SAT Solver Performance

## Graphs and their Community Structure

- Informal definition of community in a graph
  - View SAT instances as variable-incidence graphs
  - A community is sub-graph that has more internal edges than outgoing ones
  - A community structure characterizes how “well clustered” a graph is

# A note on Community Structure in Graphs

## Applications in Various Domains

- Community structure [GN03,CNM04,OL13] is used to study all kinds of complex networks.
  - Social Networks, e.g. Facebook
  - Internet
  - Protein networks
  - Neural network of the human brain
  - Citation graphs
  - Business networks
  - Populations
  - And more recently, the graph of logical formulas

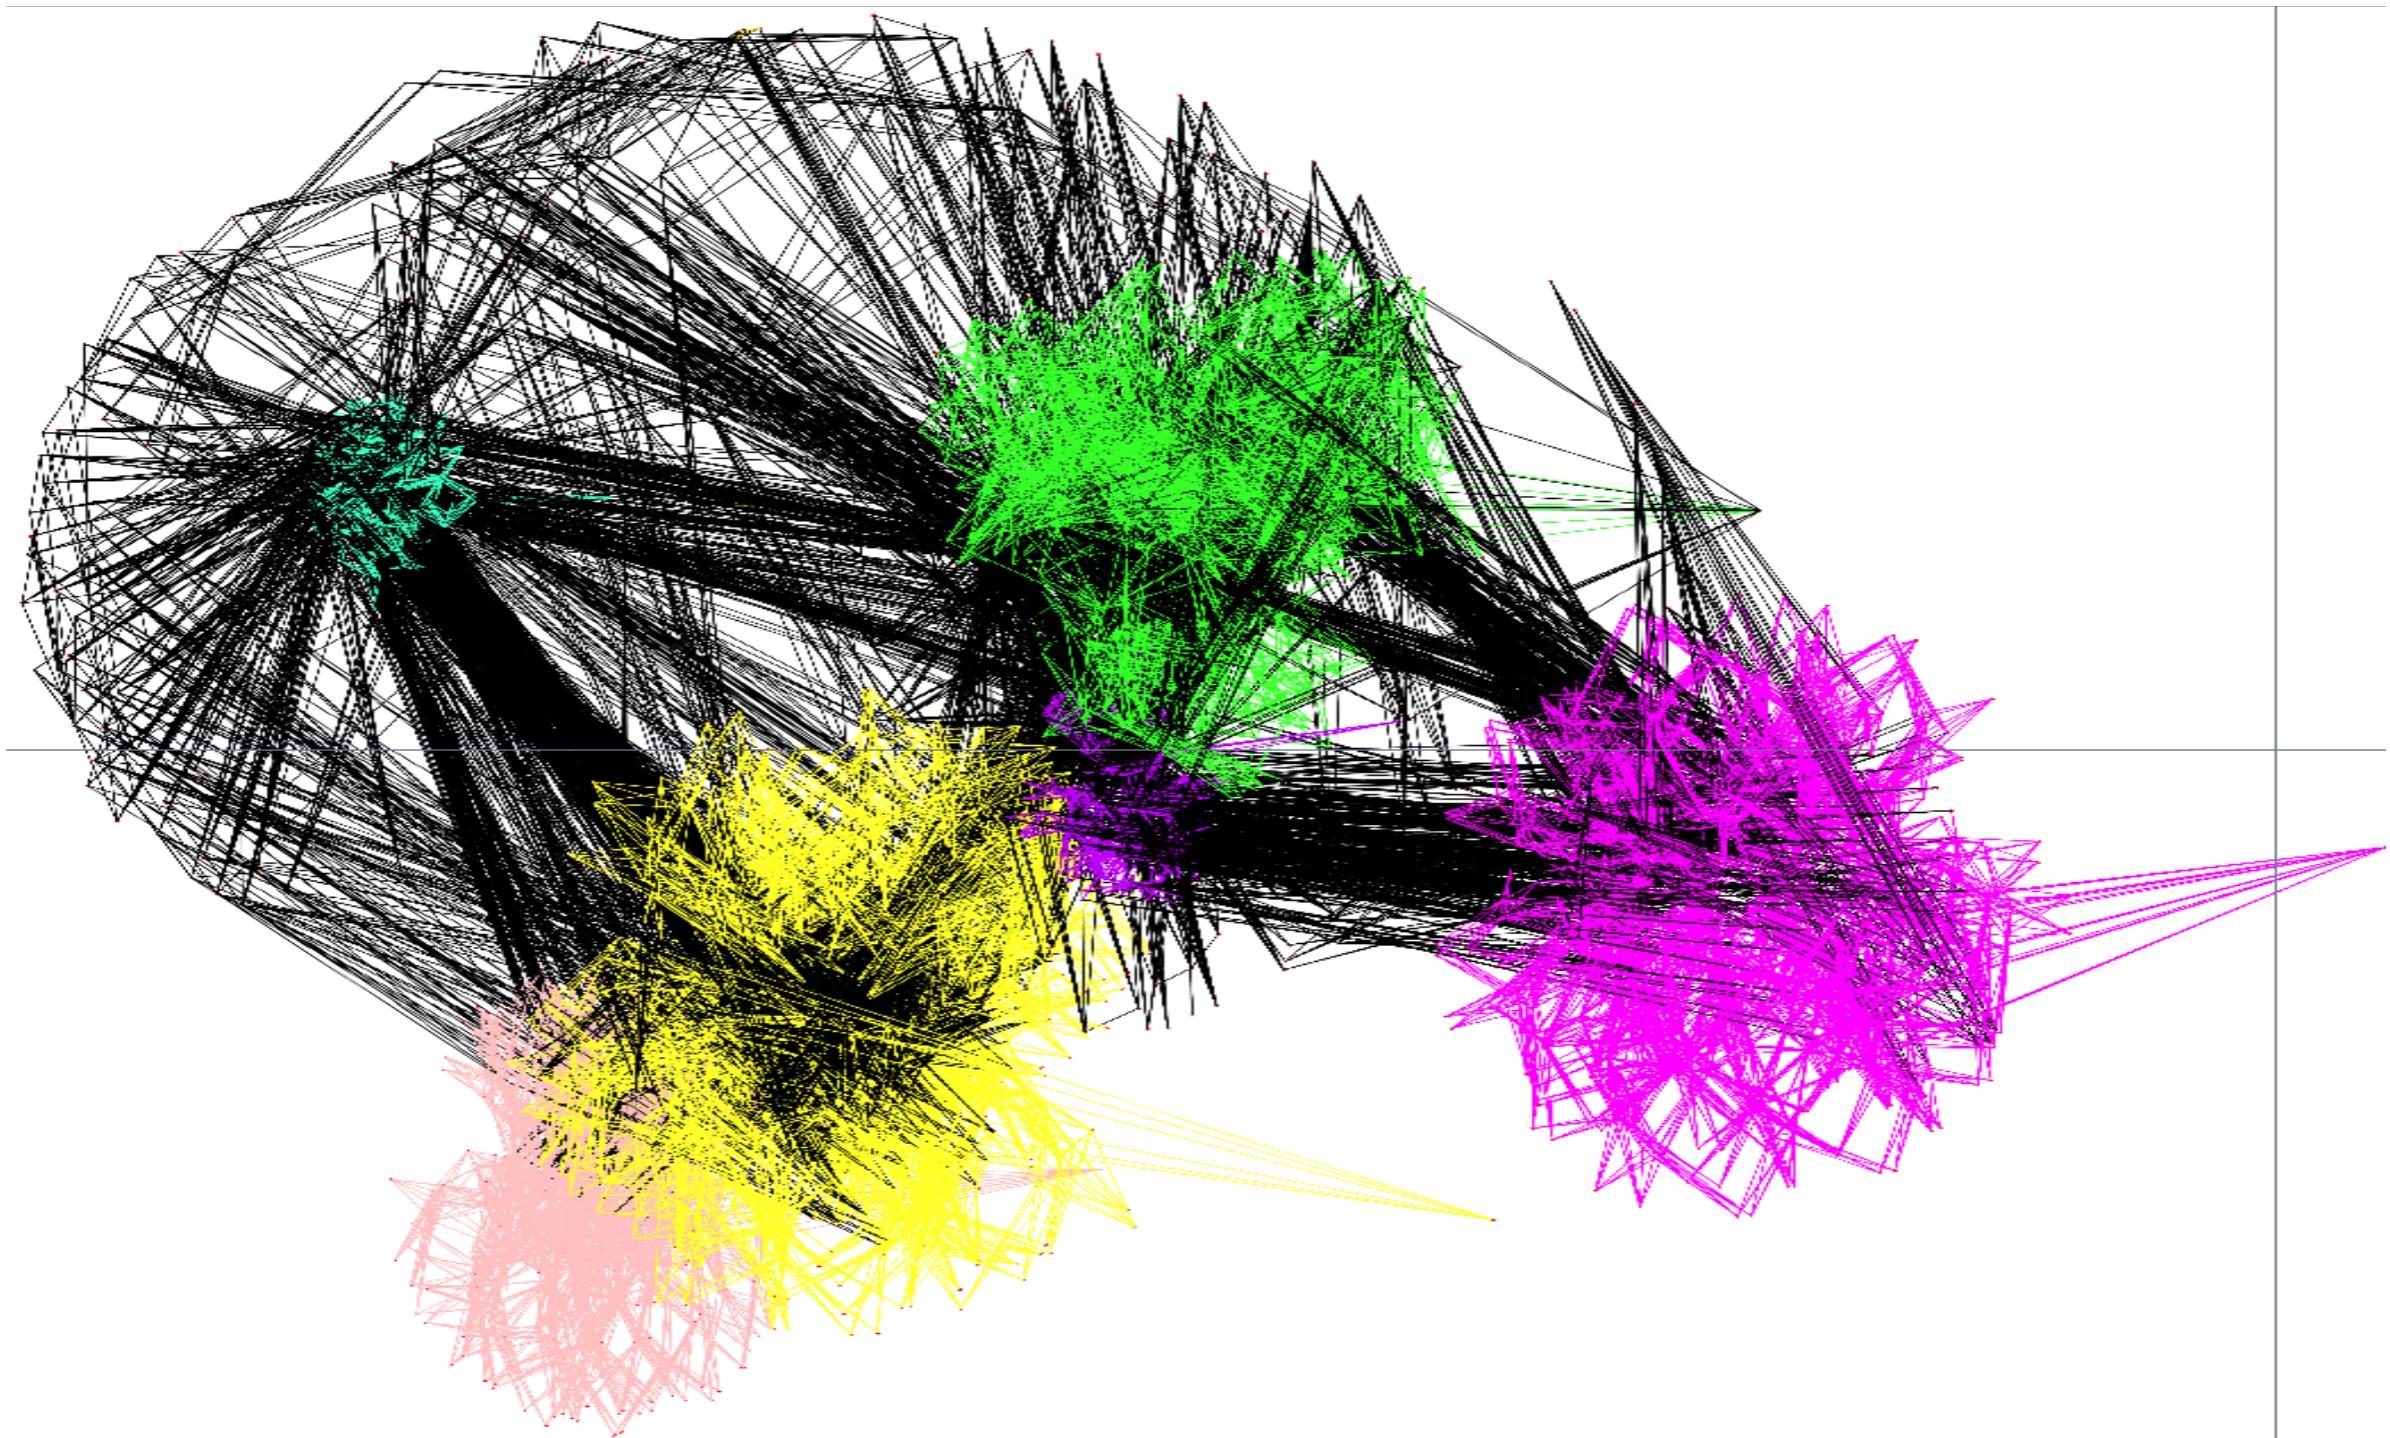
# Modularity and Communities in Graphs

## Community Structure in Graphs

- Modularity ( $Q$ ) of graph lies between 0 and 1
- $Q$  measures quality
  - How separable are the communities in the graph
  - Higher  $Q$  implies “good community structure”, i.e. *highly separable communities*
  - Lower  $Q$  implies “bad community structure”, i.e. *one giant hairy ball*

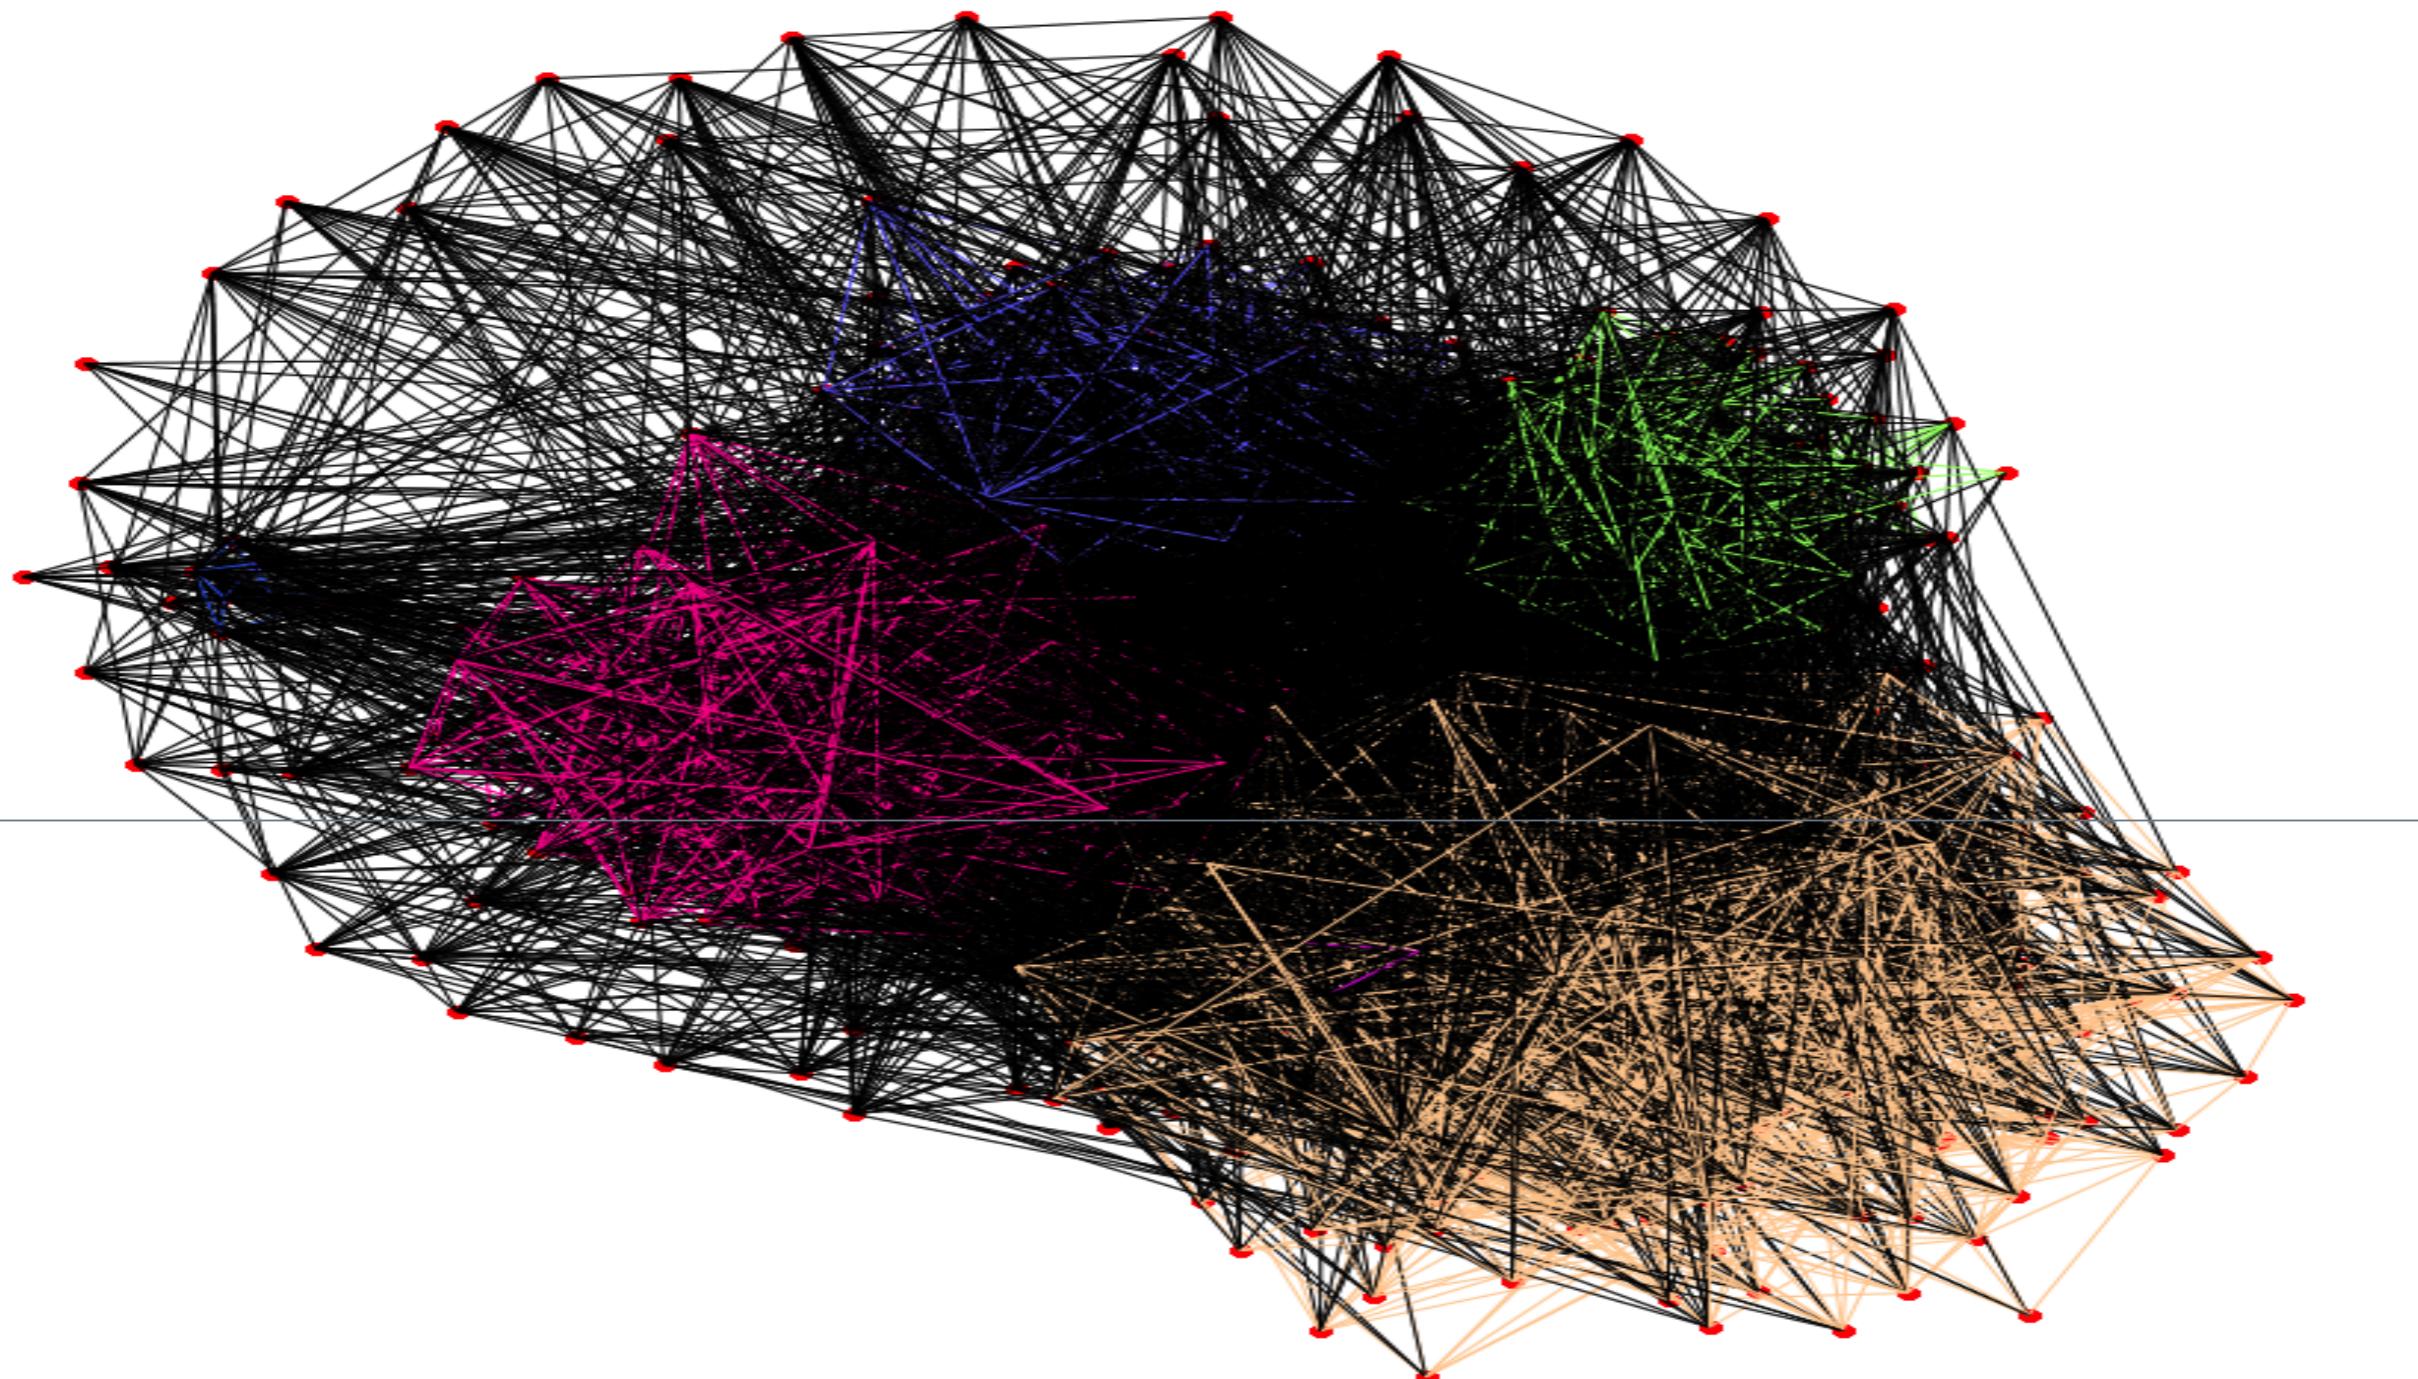
# Community Structure in Graphs

Variable-incidence Graph of Non-random Formula



# Community Structure in Graphs

Variable-incidence Graph of Random Formula



# Modularity and Communities in Graphs

## Community Structure in Graphs

- How to compute community structure?
- The decision version of the  $Q$  maximisation problem is NP-complete [Brandes et al., 2006]
- Many efficient *approximate* algorithms proposed, e.g. [CNM04] and [OLI13]
- We use the above two algorithms for our experiments
- We get similar results with both algorithms, increasing our confidence in our results

# Community Structure and Random Instances

## Experiments #1: Hypothesis and Definitions

Hypothesis tested:

- Is there a range of  $Q$  values for randomly generated instances, that are hard for CDCL solvers; regardless of the number of clauses/variables
- Are randomly generated instances outside this range uniformly easy

# Community Structure and Random Instances

## Experiments #1: Setup

- Randomly generated 550,000 SAT instances for the experiment
  - Varied  $N_V$  between 500 and 2000 in increments of 100
  - Varied  $N_{cl}$  between 2000 and 10000 in increments of 1000
  - Varied target Q between 0 and 1 in increments of 0.01
  - Varied “Number of communities” between 20 and 400 in increments of 20
- Experiments using MiniSAT
  - Timeout of 900 seconds per run
  - Run solver on inputs in a random order
  - Average the running time over several runs

# Community Structure and Random Instances

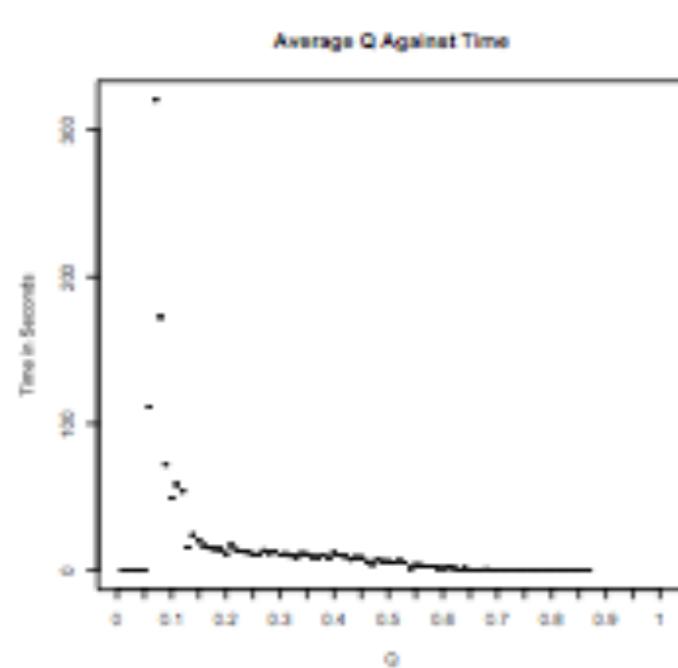
## Experiments Performed (#1)

- Plotted Q against time
- Noticed significant increase in execution time when  $0.05 \leq Q \leq 0.13$
- Also recomputed the results using a stratified sample
  - Used due to high number of instances within target range
  - Randomly sample the data taking 250 results from each 0.1 range of Q between 0 and 0.9
  - Almost the same result:  $0.05 \leq Q \leq 0.12$

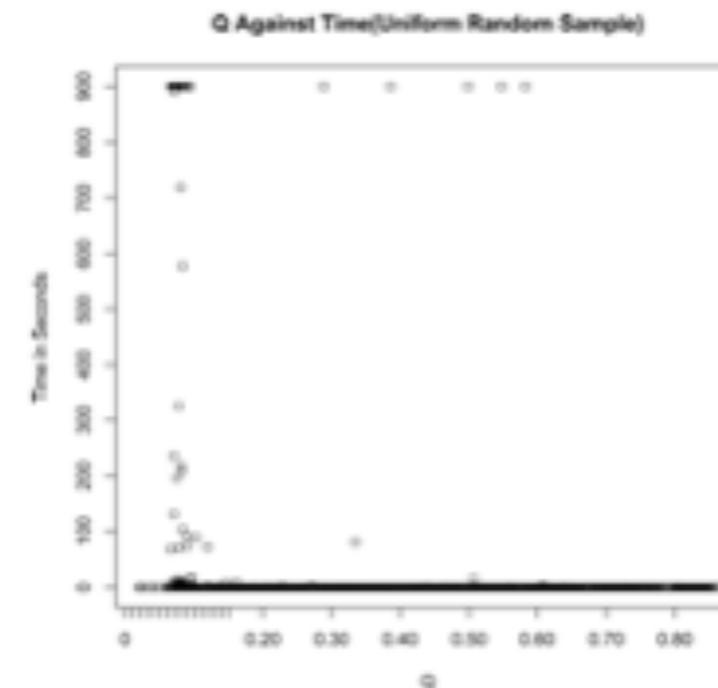
# Community Structure and Random Instances

## Experiments Performed (#1)

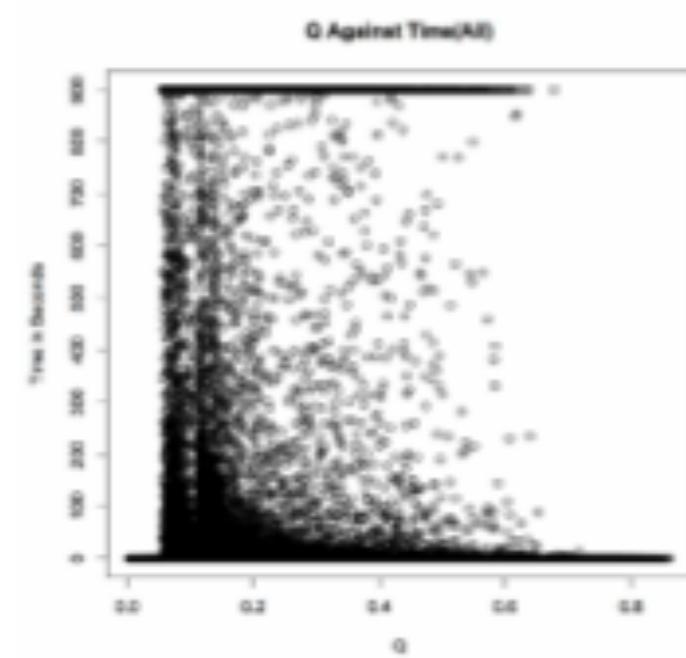
- Huge increase in running time of randomly generated instances when  $0.05 \leq Q \leq 0.13$



(a) Average Time



(b) Stratified Sample



(c) All instances

# Community Structure and Random Instances

## Experiments #2: Hypothesis and Definitions

Hypothesis tested:

- Are the community modularity and number of communities better correlated with the running time of CDCL solvers than traditional metrics
- Is the correlation better for industrial instances than randomly generated or hand crafted ones

# Community Structure and Random Instances

## Experiments #2: Hypothesis and Definitions

- Instances used
  - Approximately 800 instances from the SAT 2013 competition. For the remaining we couldn't compute community structure due to resource constraints
  - Using OL algorithm to compute community structure for the 800 instances
    - Much faster and more scalable
- All experimental results are for Minipure
  - Obtained from the SAT 2013 competition website
- Used statistical tool R to perform standard linear regression

# Community Structure and Random Instances

## Experiments Performed (#2)

- Performed linear regression on the solver running time twice
  - Once with community structure metrics (and variables/clauses)
  - Once without
- Compared the adjusted  $R^2$  (variability) from both experiments
  - Variability measures how good the models predicted results are, compared with the actual results.
  - Varies from 0 to 1
- The lower the variability (higher the  $R^2$ ) the more predictive the model

# Community Structure and Random Instances

## Experiments Performed (#2)

- Timeouts included
  - A large portion (Approximately 60%) of the instances timedout
  - Not ideal, but without them there isn't enough data
- $\log(\text{time})$  used
  - Timeouts
  - Wide distribution between instances that finished and timedout
- Data standardized to have mean = 0 and standard deviation = 1
  - Standard practice when regressors are in different scales.

# Community Structure and Random Instances

## Experiments Performed (#2)

- Model #1 -  $R^2 \sim 0.5$ 
  - $\log(\text{time}) \sim |\text{CL}| * |\mathcal{V}| * Q * |\text{CO}| * \text{QCOR} * \text{CLVR}$
  - \* denotes interaction terms between factors
  - $|\text{CL}|$  = number of clauses
  - $|\mathcal{V}|$  = number of variables
  - $|\text{CO}|$  = number of communities
  - QCOR = ratio of Q to communities
  - CLVR = ratio of clauses to variables
- Model #2 -  $R^2 \sim 0.33$ 
  - $\log(\text{time}) \sim |\text{CL}| * |\mathcal{V}| * \text{CLVR}$

# Community Structure and Random Instances

## Experiments #2: Results and Interpretation

- The regressions show us that the model with the community structure metrics is a better predictor of running time than traditional metrics, i.e. number of clauses/variables.

Factor	Estimate	Std. Error	t value	Pr(>  t )	Sig
$ CO $	-1.237e+00	3.202e-01	-3.864	0.000121	***
$ CL  \odot Q \odot QCOR$	-4.226e+02	1.207e+02	-3.500	0.000492	***
$ CL  \odot Q$	-2.137e+02	6.136e+01	-3.483	0.000523	***
$ CL  \odot Q \odot  CO  \odot QCOR \odot VCLR$	-1.177e+03	3.461e+02	-3.402	0.000702	***
$ CL  \odot Q \odot  CO $	-6.024e+02	1.774e+02	-3.396	0.000719	***
$Q \odot QCOR$	3.415e+02	1.023e+02	3.339	0.000881	***
$Q$	1.726e+02	5.200e+01	3.318	0.000947	***
$Q \odot  CO  \odot QCOR$	9.451e+02	2.927e+02	3.229	0.001292	**

# Literal Block Distance (LBD) and Communities

## Experiment #3: Hypothesis and Definitions

### Hypothesis tested

- The number of communities in a conflict clause correlates strongly with its LBD measure

### What is LBD? (First introduced in Glucose solver [AS09])

- LBD measure  $M$  of a learnt clause  $C$  is a rank based on the number  $N$  of distinct decision levels the vars in  $C$  belong to
  - The lower the value of  $N$  the better the rank  $M$
  - LBD is a powerful measure of the utility of a conflict clause

# Literal Block Distance (LBD) and Communities

## Experiment #3: Hypothesis and Definitions

- Clause deletion
  - Integral to the efficiency of modern solvers
  - Without clause deletion, conflict clause production quickly consumes available memory
- Which clauses to delete? LBD to the rescue
  - Periodically delete conflict clauses with bad LBD rank
  - As we will see, clauses with bad LBD rank are shared by many communities

# Literal Block Distance (LBD) and Communities

## Experiment #3: Intuition

- The number of communities in a conflict clause
  - The number of communities  $N$  in a conflict clause  $C$  is the number of distinct communities the variables in  $C$  belong to
- Intuition behind the hypothesis
  - High quality conflict clauses tend to span very few communities, i.e.  $N$  is small
  - High quality conflict clauses are likely to cause more propagation per decision variable, and hence are likely to have low LBD
  - LBD picks out high quality conflict clauses

# Literal Block Distance (LBD) and Communities

## Experiment #3: Setup

- Instances considered
  - 189 SAT 2013 industrial category instances out of 300
  - We were only able to compute communities for these 189
  - The rest caused memory-out errors
- Step I
  - For each of the 189 instances, compute:
    - Community structure
    - The number of communities a learnt clause spans
    - LBD of every learnt clause (only for the first 20,000 due to resource constraints)

# Literal Block Distance (LBD) and Communities

## Experiments Performed (#3)

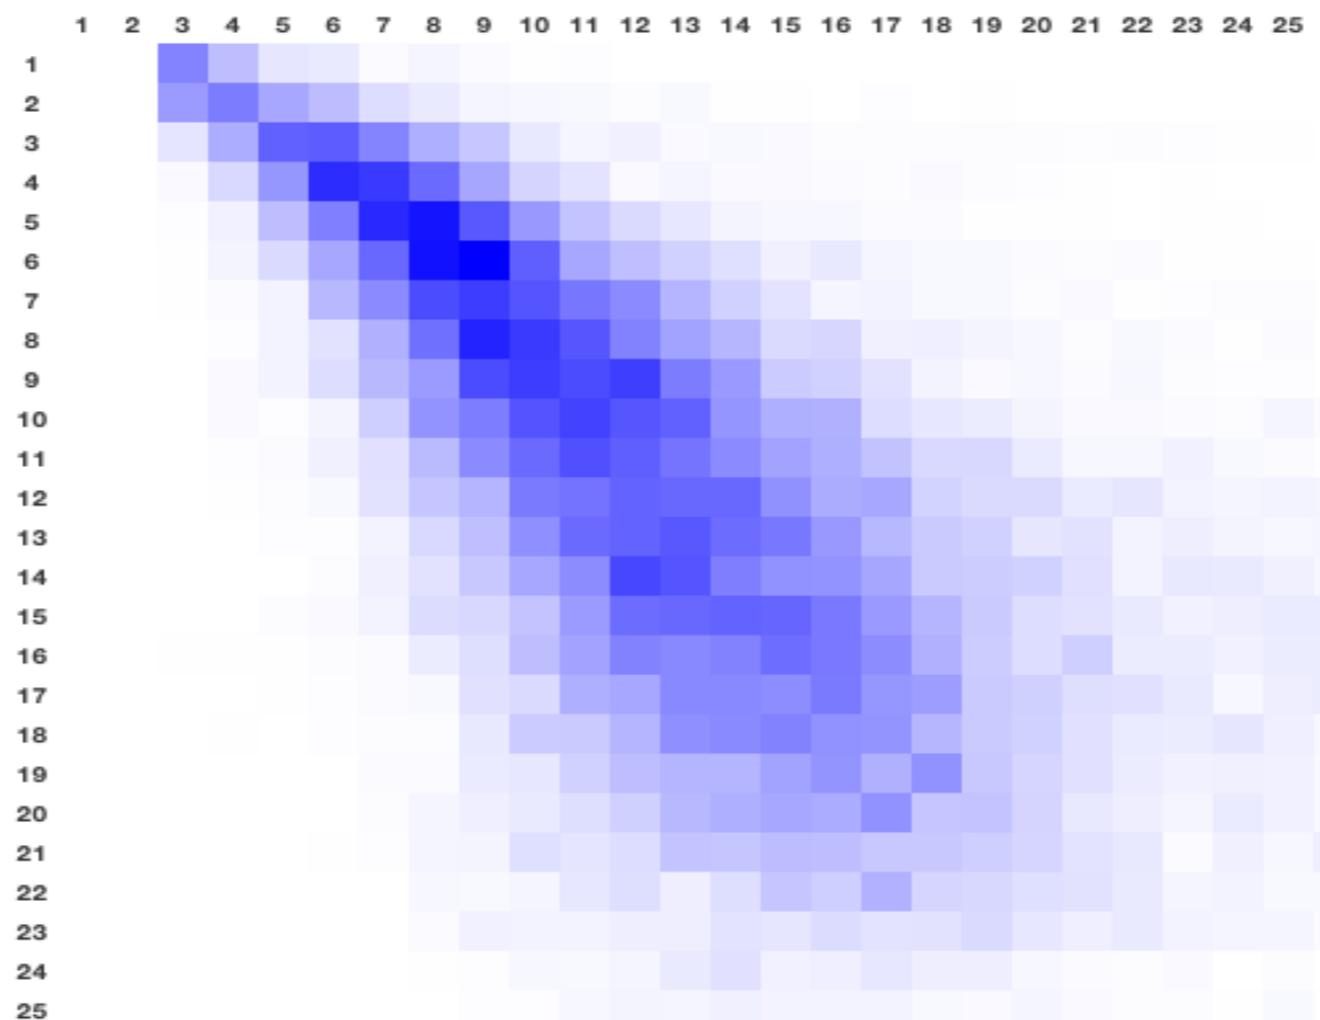
- Step 2
  - LBD of every learnt clause considered, was correlated with the number of communities it spans
  - Thousands of data points over the 189 instances
- Correlate LBD and number of communities using heatmaps
  - Heatmap of LBD and communities of learnt clauses
  - Otherwise difficult to correlate thousands of data points over hundreds of instances
  - One heatmap per SAT instance

# Literal Block Distance (LBD) and Communities

## Experiments #3: Results and Interpretation

- Result #1

- Most industrial instances have a very strong correlation between LBD and communities



# Impact of Community Structure and Solver Running Time

## Scope for Improvement

- Use larger set of data
  - Enabling us to exclude timeouts
- Focus on individual categories
  - In some cases the  $R^2$  for the individual categories is higher than the complete set
- Consider different regression techniques
  - The non-normality of the data stops us from estimating confidence intervals
- Try experiments on more solvers
  - Glucose result is similarly strong
- Compare different random generation techniques

# Impact of Community Structure and Solver Running Time

## Future Work

- Consider different graph representations for community detection
- Is it possible to determine a highly predictive model?
- Compare community structure model with graph-width based models
- Other solver measures, e.g. memory usage, rate of conflict generation
- Build visualisation tool to confirm/refute hypothesis
  - SATGraf
- Dynamically track how community structure of learnt clauses evolve

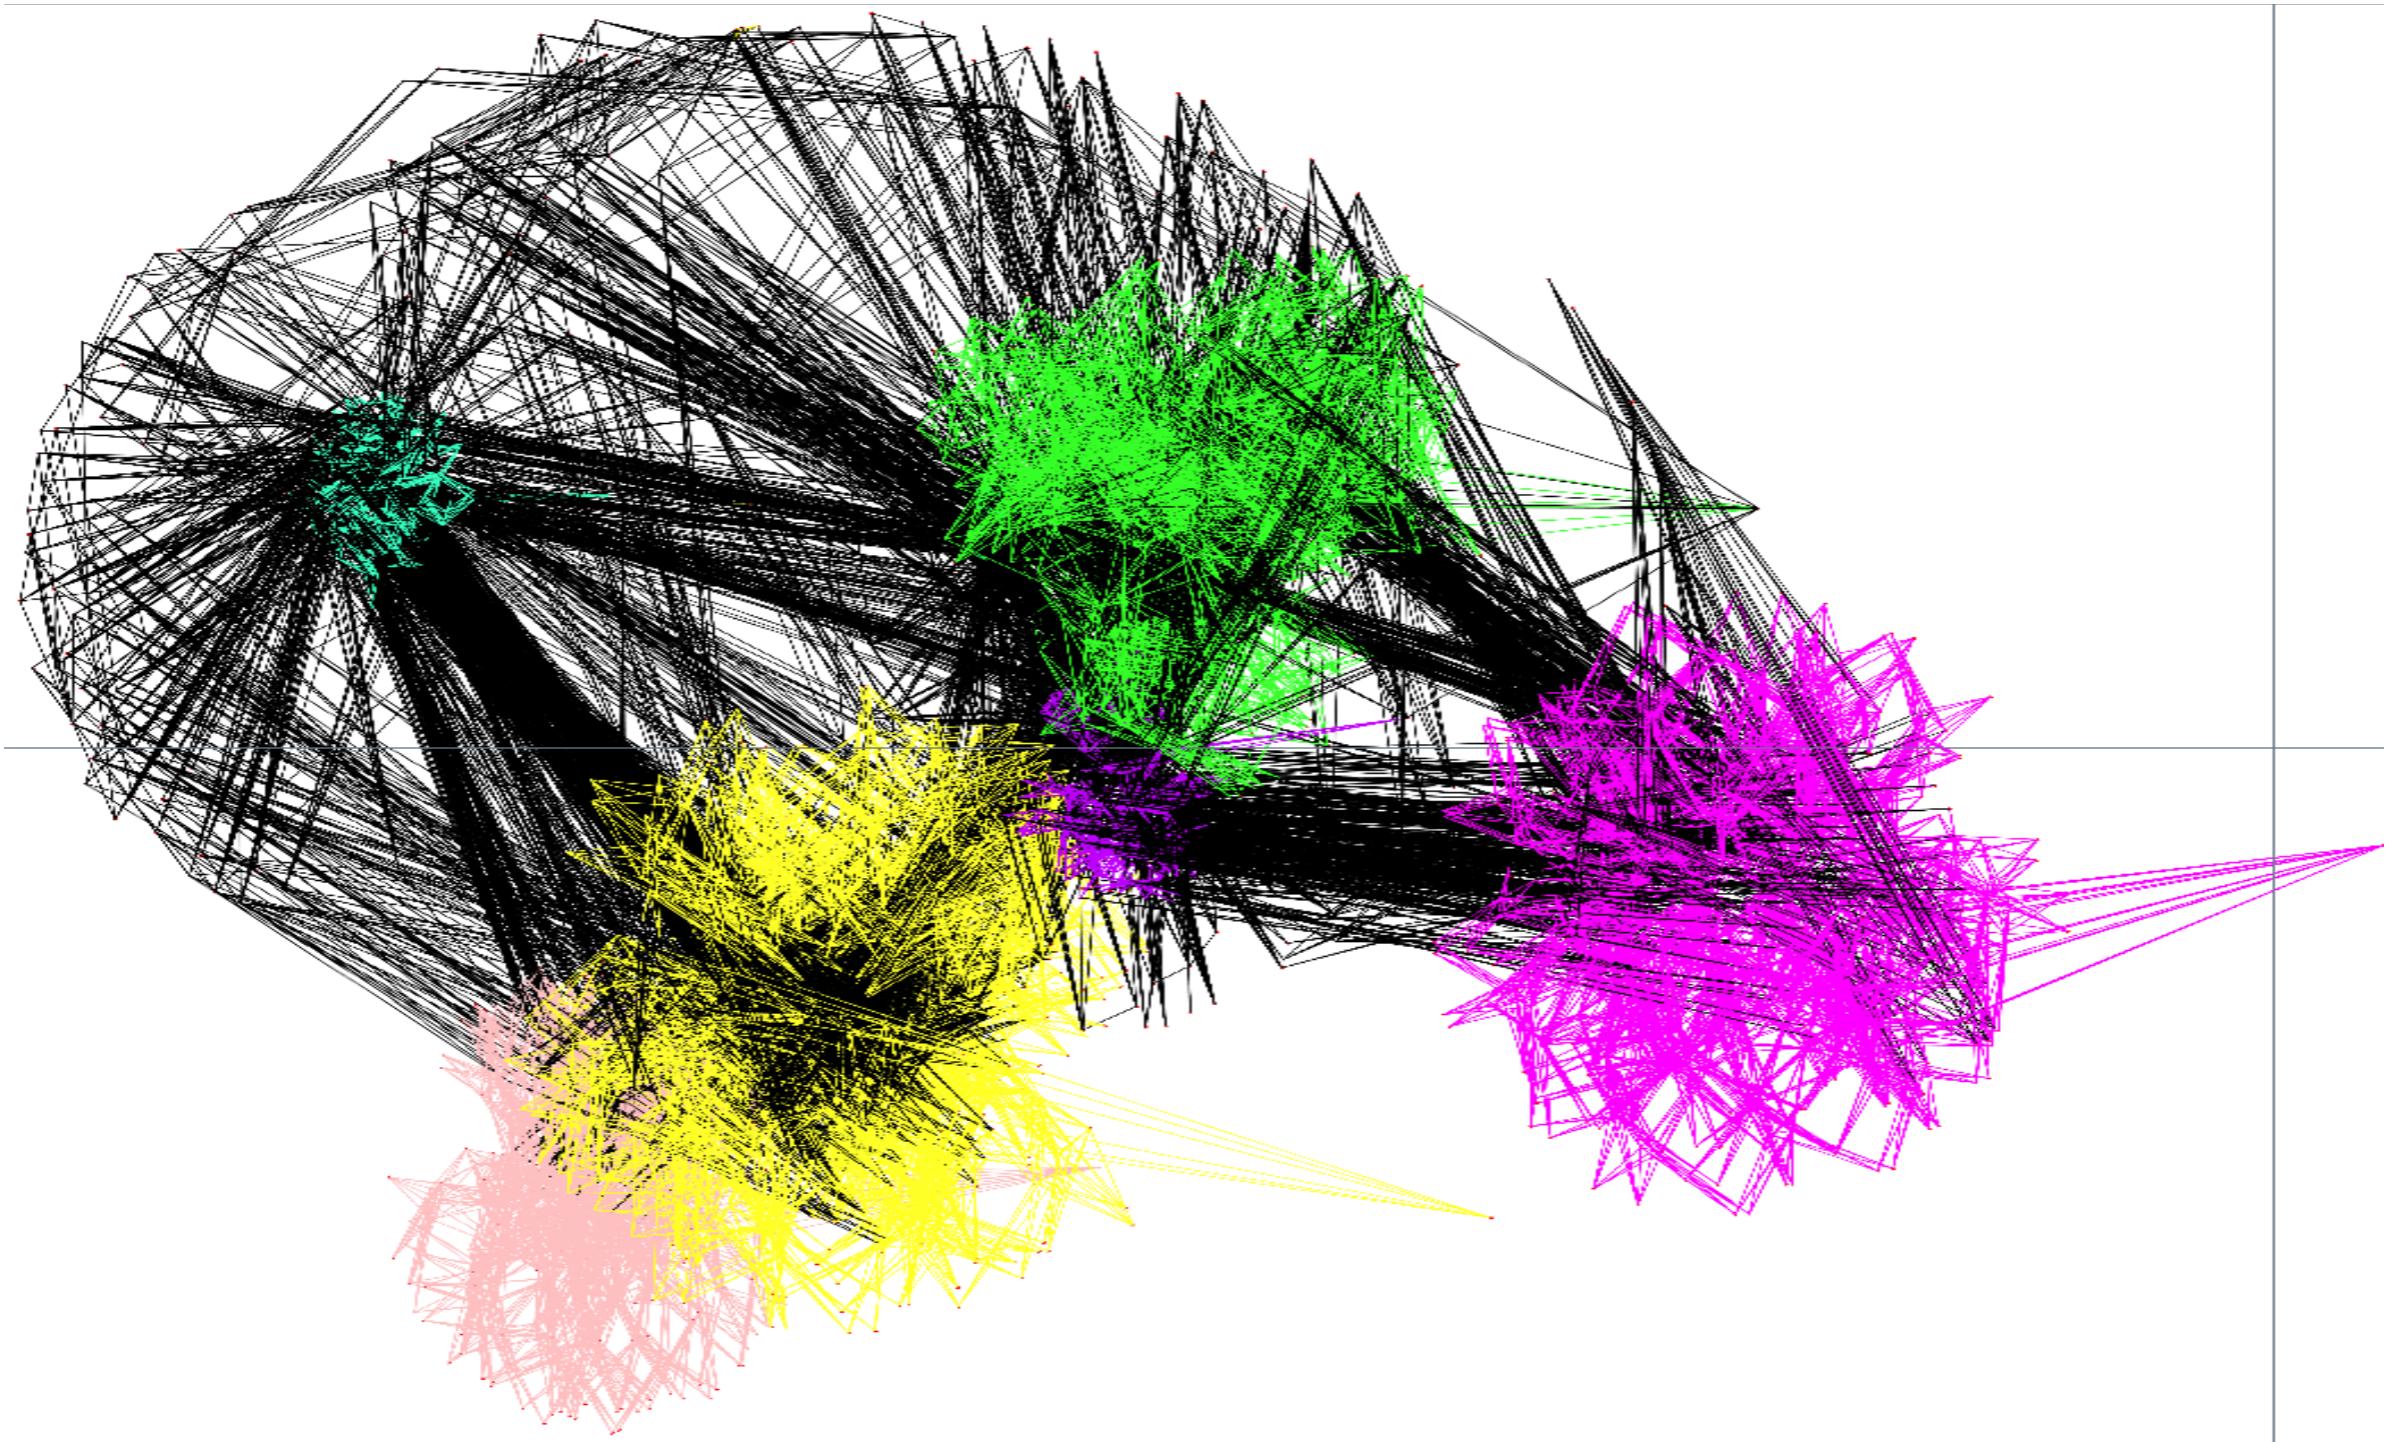
# Community Structure and SAT Solver Performance

## Conclusions

- Take home message
  - Community structure of SAT instances strongly affect solver performance
- Result #1
  - Hard random instances have low Q ( $0.05 \leq Q \leq 0.13$ )
- Result #2
  - Number of communities and Q of SAT instances are more predictive of Minipure performance than other measures
- Result #1
  - Strong correlation between community structure and LBD in Glucose

# Community Structure and SAT Solver Performance

## Questions



# Key Contributions

<http://ece.uwaterloo.ca/~vganesh>

Name	Key Concept	Impact	Pubs
<b>STP</b> Bit-vector & Array Solver <sup>1,2</sup>	Abstraction-refinement for Solving	Concolic Testing	CAV 2007 CCS 2006 TISSEC 2008
<b>HAMPI</b> String Solver <sup>1</sup>	App-driven Bounding for Solving	Analysis of Web Apps	ISSTA 2009 <sup>3</sup> TOSEM 2012 CAV 2011
<b>Z3-str</b> String and Numeric Solver <sup>1</sup>	Solver combination	Analysis of Web Apps	FSE 2013
<b>Taint-based Fuzzing</b>	Information flow is cheaper than concolic	Scales better than concolic	ICSE 2009
<b>Automatic Input Rectification</b>	Acceptability Envelope: Fix the input, not the program	New way of approaching SE	ICSE 2012
<b>Undecidability of forall-exists fragment over word equations</b>	Reduction from halting problem for two-counter machines		HVC 2012

1. 100+ research projects use STP and HAMPI
2. STP won the SMTCOMP 2006/2010 and second in 2011/2014 competitions for bit-vector solvers
3. HAMPI: ACM Best Paper Award 2009
4. Retargetable Compiler (DATE 1999, 2008). Proof-producing decision procedures (TACAS 2003)